Teaching large lan
“forget” unwantec

While large language models (LLMs) are becomin
amounts of data, a new technique that does the c
unlearning.

This relatively new approach teaches LLMs to for
copyrighted data. It is faster than retraining mod
specific unwanted data or behavior.

No surprise then that tech giants like IBM, Googl
unlearning ready for prime time. The growing foc
hiccups with this technique: models that forget to
evaluate the effectiveness of the unlearning.

From learning to unle

Trained on terabytes of data, LLMs "learn" to mal
explicitly programmed to do so. This branch of AI
popularity as machine learning algorithms imitat
improving the accuracy of the content they gener

But more data also means more problems. Or as
Baracaldo puts it, "Whatever data is learned—the

And so ever larger models can also generate mor
data that defy cybersecurity standards. Why? The
untrusted data from the internet. Even with rigor
define what questions *not* to answer and what an
inspect a model's output—still, unwanted behavi
creep through.

Retraining these models to remove the undesirab
dollars. Furthermore, when models are open sou
carried onto many other models and applications

Unlearning approaches aim to alleviate these pro
specific data points like content containing harm
unwanted text prompts, unlearning algorithms ef

Mixture of Experts | 14 November, episode

Decoding AI: Weekly News Rou

Join our world-class panel of engineers,
as they cut through the AI noise to bring

Watch all episodes of Mixture of Experts →

Forgetting Harry Pott

A team of researchers from Microsoft used this u
Meta's Llama2-7b model forget copyrighted mate
on from the internet. Before unlearning, when the
Harry Potter?" the model responded: "Harry Pott
series of fantasy novels."

After fine-tuning the model to "unlearn" copyrigh
following to the same prompt: "Harry Potter is a I

"In essence, every time the model encounters a c
original content," explained the researchers Rone
The team shared their model on Hugging Face
and tinker with it as well.

In addition to removing copyrighted material, ren
privacy is another high-stake use case. A team, le
at Austin, collaborating with AI specialists at JP M
for image-to-image generative models. In a recen
eliminate unwanted elements of images (the "for
the overall image set.

This technique could be helpful in scenarios such
instance, said Professor Marculescu. "If there we
those out to protect their privacy."

Google is also busy tackling unlearning within the
June 2023, Google launched its first machine unl
an age predictor that had been trained on face im
training images had to be forgotten to protect the

While it's not perfect, the early results from vario
unlearning on a Llama model, for instance, Barac
toxicity score from 15.4% toxicity to 4.8% withou
performed. And instead of taking months to retra
took all of 224 seconds.

Speed bumps

So why isn't machine unlearning widely used?

"Methods to unlearn are still in their infancy and

The first challenge that looms large is "catastropl
than the researchers wanted so the model then r

The IBM team has developed a new framework t
training. Using an approach they describe as split
able to unlearn undesirable behavior such as toxi
biosecurity or cybersecurity risks, while preservir

Developing comprehensive and reliable evaluatic
unlearning efforts also remains an issue to resolv

The future of machin

While unlearning may still be finding its feet, rese
broad array of potential applications, industries a

In Europe for instance, the EU's General Data Pro
be forgotten." If an individual opts to remove thei
companies adhere to this legislation and remove critical data. Beyond security and privacy, machine unlearning could also be useful in any situation where data needs to be added or removed when licenses expire or clients, for instance, leave a large financial institution or hospital consortium.

"What I love about unlearning," says Baracaldo, "Is that we can keep using all our other lines of defense like filtering data. But we can also 'patch' or modify the model whenever we see something goes wrong to remove everything that's unwanted."

Author

Aili McConnon
Staff Writer
IBM

Report

IBM is named a Leader in Data Science & Machine Learning

Learn why IBM has been recognized as a Leader in the 2025 Gartner® Magic Quadrant™ for Data Science and Machine Learning Platforms.

Read the report →

Resources

Report

The 2025 CEO's guide: 5 mindshifts to supercharge business growth

Activate these five mindshifts to cut through the uncertainty, spur business reinvention, and supercharge growth with agentic AI.

Read the report →

Training

Level up your ML expertise

Learn fundamental concepts and build your skills with hands-on labs, courses, guided projects, trials and more.

Explore ML courses

Ebook

Unlock the power of generative AI + ML

Learn how to confidently incorporate generative AI and machine learning into your business.

Read the ebook →

Guide

Put AI to work: Driving ROI with gen AI

Want to get a better return on your AI investments? Learn how scaling gen AI in key areas drives change by helping your best minds build and deliver innovative new solutions.

Read the guide →

Ebook

How to choose the right

AI models

Explore IBM Granite

IBM® Granite™ is our family of open, performant and trusted AI models, tailored for business and optimized to scale your AI applications. Explore language, code, time series and guardrail options.

Meet Granite →

Guide

How to thrive in this new era of AI with trust and confidence

Dive into the 3 critical elements of a strong AI strategy: creating a competitive edge, scaling AI across the business and advancing trustworthy AI.

Read the guide

Report

AI in Action Report

We surveyed 2,000 organizations about their AI initiatives to discover what’s working, what’s not and how you can get ahead.

1 / 2

Related solutions

IBM® watsonx Orchestrate™
Easily design scalable AI assistants and agents, automate repetitive tasks and simplify complex processes with IBM® watsonx Orchestrate™.

Explore watsonx Orchestrate →

AI for developers
Move your applications from prototype to production with the help of our AI development solutions.

Explore AI development tools →

AI consulting and services
Reinvent critical workflows and operations by adding AI to maximize experiences, real-time decision-making and business value.

Explore AI services →

Take the next step

Whether you choose to customize pre-built apps and skills or build and deploy custom agentic services using an AI studio, the IBM watsonx platform has you covered.

Explore watsonx Orchestrate

Explore watsonx.ai