



Large Language Models for Tabular Data: Progresses and Future Directions

Haoyu Dong*
Microsoft AI
Beijing, China
haoyu.dong@microsoft.com

Zhiruo Wang
Carnegie Mellon University
Pittsburgh, USA
zhiruow@cs.cmu.edu

ABSTRACT

Tables contain a significant portion of the world's structured information. The ability to efficiently and accurately understand, process, reason about, analyze, and generate tabular data is critical for achieving Artificial General Intelligence (AGI) systems.

However, despite their prevalence and importance, tables present unique challenges due to their structured nature and the diverse semantics embedded within them. Textual content, numerical values, visual formats, and even formulas in tables carry rich semantic information that is often underutilized due to the complexity of accurately interpreting and integrating.

Fortunately, the advent of Large Language Models (LLMs) has opened new frontiers in natural language processing (NLP) and machine learning (ML), showing remarkable success in understanding and generating text, code, etc. Applying these advanced models to the domain of tabular data holds the promise of significant breakthroughs in how we process and leverage structured information.

Therefore, this tutorial aims to provide a comprehensive study of the advances, challenges, and opportunities in leveraging cutting-edge LLMs for tabular data. By introducing methods of prompting or training cutting-edge LLMs for table interpreting, processing, reasoning, analytics, and generation, we aim to equip researchers and practitioners with the knowledge and tools needed to unlock the full potential of LLMs for tabular data in their domains.

CCS CONCEPTS

• **Information systems** → **Information retrieval**.

KEYWORDS

Tabular data, Large language models, Representation learning

ACM Reference Format:

Haoyu Dong and Zhiruo Wang. 2024. Large Language Models for Tabular Data: Progresses and Future Directions. In *Proceedings of the 47th International ACM SIGIR Conference on Research and Development in Information Retrieval (SIGIR '24)*, July 14–18, 2024, Washington, DC, USA. ACM, New York, NY, USA, 4 pages. <https://doi.org/10.1145/3626772.3661384>

*Corresponding author.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

SIGIR '24, July 14–18, 2024, Washington, DC, USA.

© 2024 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM ISBN 979-8-4007-0431-4/24/07

<https://doi.org/10.1145/3626772.3661384>

1 DETAILED SCHEDULE

The tutorial, as an in-person event, adopts the length of three hours plus breaks. A detailed schedule is as follows:

- **PART I.** Introduction to This Tutorial (20 minutes)
 - Motivation (10 minutes)
 - Scope (i.e. breadth/depth) (10 minutes)
- **PART II.** Table Representation for LLMs (35 minutes)
 - Fundamental elements of tabular data (5 minutes)
 - Representation approaches of tabular data (20 minutes)
 - Retrieval or compression methods for large tables or databases (10 minutes)
- **PART III.** Training LLMs for Tabular Data (25 minutes)
 - Corpora of tabular data (5 minutes)
 - Continued pre-training (10 minutes)
 - Single-task & multi-task fine-tuning (10 minutes)
 - Question and Answer Session (10 minutes)
 - Coffee Break —
- **PART IV.** Tasks and Benchmarks (20 minutes)
 - Table interpreting and processing (10 mins)
 - Table reasoning and analysis (10 mins)
- **PART V.** LLM-driven Table Agents (40 minutes)
 - Tabular environments (10 minutes)
 - Agents: decompose tasks, use tools, sampling and refine (30 minutes)
- **PART VI.** Hands-on Session and New horizons (15 minutes)
 - Demonstration & Hands-on session (10 mins)
 - New horizons (5 minutes)
 - Question and Answer Session (10 minutes)

2 RELEVANCE

This tutorial holds strong relevance to the core domains of SIGIR where the topics of table data retrieval, machine learning, and data intelligence are of paramount interest. This tutorial will not only build upon the knowledge shared in previous tutorials such as "Web Table Extraction, Retrieval and Augmentation" [67] at SIGIR 2019 and "From Tables to Knowledge: Recent Advances in Table Understanding" [47] at KDD 2021, but will also introduce the latest developments in LLMs that were not covered previously. Recent surveys on LMs for tabular data include [14, 15, 21, 41, 70].

3 BRIEF CONTENT

3.1 Structured, Semi-structured, and Image Tables

Well-structured tables, such as those found in databases, feature a clearly defined relational schema. This format facilitates precise data retrieval using formal languages like SQL, R, and Python's

Pandas [15]. In contrast, semi-structured tables, common in web pages, LaTeX documents, spreadsheets, and Word documents, often lack necessary meta-information to accurately capture the table's high-level structure and data types, posing challenges for precise and automatic information retrieval [17, 35]. Additionally, a substantial volume of image tables exists, containing only raw RGB data [14], and PDF tables are typically treated as such [49].

3.2 Table Representation for LLMs

The process of converting tables into prompts while preserving their semantic integrity presents the first challenge. LLMs necessitate prompts in a linear, sequential text format, which contrasts with the inherently two-dimensional structure of tables that possess a defined schema. The prompt might encompass its content, schema, visual formats, foreign keys, formulas, etc.

A prevalent method to encode table content involves linearizing the tables row by row, employing column separators such as “|” to demarcate individual cells. Studies have assessed LLMs’ effectiveness across a spectrum of table representations, including CSV, markdown, JSON, XML, and HTML [50, 51].

Studies found that incorporating details like foreign keys and minimizing explanations can notably improve LLM performance on these tasks. To establish database connections, DBCopilot [53] streamlines the text-to-SQL conversion process by segmenting it into schema routing and SQL generation phases.

Visual formats play a vital role in interpreting tables, and studies exist on LLMs’ use of image-based table representations [14].

3.3 Retrieval or Compression Methods for Large Tables or Databases

Even LLMs with extensive contextual capabilities may fall short when dealing with exceedingly large tables. Furthermore, [6] discovered that GPT’s effectiveness dwindles to mere conjecture with tables exceeding 1,000 tokens in size. Additionally, longer prompts lead to increased costs, particularly for applications relying on LLM APIs.

To bridge this gap, [28, 39] suggested simple truncation strategies, limiting the input to a maximum sequence length. [51] implemented predefined constraints to align with LLM call requirements. Retrieval-augmented LLMs employ a retriever-reader framework to streamline the analysis of vast table collections by focusing on a select subset. Notable efforts [37, 38] showcase the segmentation of large tables for efficient retriever-reader training, enhancing performance. DB-GPT [60] provides a comprehensive solution that includes retrieval-enhanced generation and agent interactions, catering to industrial needs. [53] advances database querying by segmenting the process into schema understanding and SQL generation, thereby offering improved accuracy and scalability for handling extensive databases across domains.

3.4 Training LLMs for Tabular Data

Before the appearance of LLMs, researchers have introduced various Tabular Language Models (TaLMs) as summarized by [15]. These models modify LMs’ structures, encoding, and objectives to align with table-centric tasks. And recent works on LLM for tabular data commonly choose not maintain the model structure.

For multi-task fine-tuning, instruction tuning transcends the limitations of fine-tuning on singular datasets by embracing the fine-tuning of language models across a spectrum of datasets. This approach significantly bolsters performance on novel tasks and enhances model flexibility [58]. Table-GPT, for instance, employs a synthesis-followed-by-augmentation strategy for dataset creation, thereby facilitating a deep understanding of tabular data structures and relationships [36]. Projects like TableLlama [68] prefer real-world data from established datasets to synthesized data, showcasing the generalized applicability of LLMs post-instruction tuning. TableLLM¹, a 13B LLM, is designed for tabular data tasks in documents and spreadsheets, suitable for Office environments. The training employs a distant supervision approach, incorporating a reasoning process extension and cross-validation to ensure data quality. The Tabular LLM also aims to gather datasets for tasks like table QA and text generation, converting them into instruction-tuning formats² to enhance LLMs’ understanding of tabular data.

Multi-modal training is essential for handling image-based tables or leveraging visual formats like human-beings. While [14] provides insights into the strengths and limitations of LLMs in processing multimodal tabular data representations, this area remains underexplored and holds significant potential for future research.

3.5 Table Tasks and Benchmarks

Table interpreting and processing. Starting with the most basic process of table collection, many earlier works collect spreadsheet [10, 19] and HTML [13, 34] tables from the web. Relatedly, the tasks of table detection and classification are essential to building a table corpus of high quality. Table datasets with table type annotations such as FUSE [2], ENRON [26], EUSES [23], and DECO [32] can enable better table recognition systems. Besides understanding, Dong et al. [16] and Kardas et al. [31] also perform structure extraction for tables from spreadsheet files and scientific papers.

More granular than the table-level processing tasks, many tasks target detection and classification tasks on the header (region) [20] and cell levels, the latter of which is more widely explored by WebSheet [16], DeEx [32], SAUS [24], and CIUS [24]. In addition, other tasks are oriented towards knowledge in tables or concerning web data, such as column type prediction [4, 5], entity linking [3], relation identification [18, 22], table retrieval [27, 56, 66].

Table reasoning and analysis. Table reasoning tasks are mainly about verifying facts in and answering questions about tables. For fact verification, available datasets such as FreebaseQA [30] and TabFact [8] enable more accurate methods to detect factual errors [1, 25, 33, 48, 65, 72]. For table-based question answering, the WikiTableQuestions dataset [44] contains compositional questions about wiki tables, and OTT-QA [7] further extends to questions about both table and text; Spider [63] focus on database queries, HiTab [11] focus on complex hierarchical spreadsheets tables, and TabMWP [40] with math world problems ground on tables. Lastly, the data-to-text generation task assesses models in table analysis, such as ToTTo [43] on web tables and HiTab [11] on spreadsheets.

¹<https://tablellm.github.io/>

²https://github.com/PhoebusSi/Alpaca-CoT/tree/tabular_llm

3.6 LLM-driven Table Agents

LLMs face challenges with table-related tasks, such as arithmetic operations [11, 44] and factual lookups [9], due to their inability to interpret compositional queries in structured table environments. To address these limitations, LLM-powered agents emerge as a solution, capable of exerting actions on tables via various means.

Environment. Tabular information poses numerous challenges depending on the data source and the presented interface. As shown by the tasks discussed above, tables could come from various sources hence bringing along individual challenges: Wikipedia tables with semi-structured format [45], spreadsheet files with complex, hierarchical table structures [11], and even large databases that may exceed model input limit [63]. Further, depending on the table and method used, a table can be presented in different ways, such as a flattened text sequences, a structure-preserving pandas dataframe, or even a clickable UI interface supporting font and color changes. These interaction modes can be realized using certain agents with certain actions, which we introduce below.

Agent. Existing table agents are designed with techniques to decompose complex tasks, leverage external tools [55, 64], and perform sampling and iterative refinement [42]. Drawing inspiration from methods such as Chain-of-Thought (CoT) [59] and least-to-most prompting, researchers have guided LLMs to break down intricate table tasks into simpler, manageable sub-tasks [46, 57]. Similarly, for handling voluminous web tables, Dater [62] utilize LLMs to divide extensive data into smaller, more digestible pieces, then solve individual pieces to finally solve the complex query.

However, agents may not solve the tasks perfectly in one try. Therefore, research initiatives such as ReAct [61] propose an iterative, closed-loop system for table agents, which repeatedly analyze table data and user intentions, generate and execute action plans, and update their action trajectories until the desired outcome is achieved. To tackle multiple tasks at a time, StructGPT [29], adeptly addresses both table QA and text-to-SQL tasks by employing versatile actions suitable for web tables, databases, and knowledge graphs. Further, TAP4LLM [52] extends the utility of LLM agents to table augmentation tasks, showcasing the adaptability of LLM agents in more table domains. Another direction for agents to explore more solutions is sampling multiple generations and selecting the best via self-consistency [54], based on the assumption that the most common output is more likely to be correct. Nonetheless, the approach is shown effective in many works [35, 46, 62, 71].

Actions. Actions are essential to the performance of agents, which should be functionally sufficient to solve the table task, while not too difficult for the agent to acquire. Actions can be categorized into agent-inherent actions and application-specific tools. Besides the most atomic actions such as add, different works introduce application-specific actions. For example, SheetCopilot [35] add spreadsheet APIs as actions, Binder [12] use a knowledge-query action for web tables, Data-Copilot [69] add human-designed Python functions for individual downstream domains. TroVE [55] use LLMs to automatically build macro actions that suits the target task.

REFERENCES

[1] Rami Aly, Zhijiang Guo, Michael Sejr Schlichtkrull, James Thorne, Andreas Vlachos, Christos Christodoulopoulos, Oana Cocarascu, and Arpit Mittal. 2021. The

Fact Extraction and VERification Over Unstructured and Structured information (FEVEROUS) Shared Task. In *Proceedings of the Fourth Workshop on Fact Extraction and VERification (FEVER)*, Rami Aly, Christos Christodoulopoulos, Oana Cocarascu, Zhijiang Guo, Arpit Mittal, Michael Schlichtkrull, James Thorne, and Andreas Vlachos (Eds.). Association for Computational Linguistics, Dominican Republic, 1–13. <https://doi.org/10.18653/v1/2021.fever-1.1>

[2] Titus Barik, Kevin Lubick, Justin Smith, John Slankas, and Emerson Murphy-Hill. 2015. Fuse: a reproducible, extendable, internet-scale corpus of spreadsheets. In *2015 IEEE/ACM 12th Working Conference on Mining Software Repositories*. IEEE, 486–489.

[3] Chandra Sekhar Bhagavatula, Thanapon Noraset, and Doug Downey. 2015. TabEL: entity linking in web tables. In *International Semantic Web Conference*. Springer.

[4] J Chen, I Horrocks, E Jimenez-Ruiz, and C Sutton. 2019. Learning Semantic Annotations for Tabular Data. *IJCAI 2019* (2019), 2088–2094.

[5] Jiaoyan Chen, Ernesto Jiménez-Ruiz, Ian Horrocks, and Charles Sutton. 2019. Colnet: Embedding the semantics of web tables for column type prediction. In *Proceedings of the AAAI Conference on Artificial Intelligence*, Vol. 33, 29–36.

[6] Wenhui Chen. 2023. Large Language Models are few(1)-shot Table Reasoners. In *Findings of the Association for Computational Linguistics: EACL 2023, Dubrovnik, Croatia, May 2–6, 2023*, Andreas Vlachos and Isabelle Augenstein (Eds.). Association for Computational Linguistics, 1090–1100. <https://doi.org/10.18653/V1/2023.FINDINGS-EACL.83>

[7] Wenhui Chen, Ming-Wei Chang, Eva Schlinger, William Wang, and William W. Cohen. 2021. Open Question Answering over Tables and Text. *arXiv:2010.10439 [cs.CL]*

[8] Wenhui Chen, Hongmin Wang, Jianshu Chen, Yunkai Zhang, Hong Wang, Shiyang Li, Xiyu Zhou, and William Yang Wang. 2019. TabFact: A large-scale dataset for table-based fact verification. *arXiv preprint arXiv:1909.02164* (2019).

[9] Wenhui Chen, Hongmin Wang, Jianshu Chen, Yunkai Zhang, Hong Wang, Shiyang Li, Xiyu Zhou, and William Yang Wang. 2020. TabFact: A Large-scale Dataset for Table-based Fact Verification. In *Proc. of ICLR*.

[10] Zhe Chen and Michael Cafarella. 2013. Automatic web spreadsheet data extraction. In *Proceedings of the 3rd International Workshop on Semantic Search over the Web*. 1–8.

[11] Zhoujun Cheng, Haoyu Dong, Zhiruo Wang, Ran Jia, Jiaqi Guo, Yan Gao, Shi Han, Jian-Guang Lou, and Dongmei Zhang. 2022. HiTab: A Hierarchical Table Dataset for Question Answering and Natural Language Generation. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Smaranda Muresan, Preslav Nakov, and Aline Villavicencio (Eds.). Association for Computational Linguistics, Dublin, Ireland, 1094–1110. <https://doi.org/10.18653/v1/2022.acl-long.78>

[12] Zhoujun Cheng, Tianbao Xie, Peng Shi, Chengzu Li, Rahul Nadkarni, Yushi Hu, et al. 2023. Binding Language Models in Symbolic Languages. In *ICLR*.

[13] Eric Crestan and Patrick Pantel. 2011. Web-scale table census and classification. In *Proceedings of international conference on Web search and data mining*.

[14] Naihao Deng, Zhenjie Sun, Ruiqi He, Aman Sikka, Yulong Chen, Lin Ma, Yue Zhang, and Rada Mihalcea. 2024. Tables as Images? Exploring the Strengths and Limitations of LLMs on Multimodal Representations of Tabular Data. *arXiv preprint arXiv:2402.12424* (2024).

[15] Haoyu Dong, Zhoujun Cheng, Xinyi He, Mengyu Zhou, Anda Zhou, Fan Zhou, Ao Liu, Shi Han, and Dongmei Zhang. 2022. Table pre-training: A survey on model architectures, pre-training objectives, and downstream tasks. *arXiv preprint arXiv:2201.09745* (2022).

[16] Haoyu Dong, Shijie Liu, Zhouyu Fu, Shi Han, and Dongmei Zhang. 2019. Semantic Structure Extraction for Spreadsheet Tables with a Multi-task Learning Architecture. In *Workshop on Document Intelligence at NeurIPS 2019*.

[17] Haoyu Dong, Shijie Liu, Shi Han, Zhouyu Fu, and Dongmei Zhang. 2019. Tablesense: Spreadsheet table detection with convolutional neural networks. In *Proceedings of the AAAI conference on artificial intelligence*, Vol. 33, 69–76.

[18] Wensheng Dou, Shi Han, Liang Xu, Dongmei Zhang, and Jun Wei. 2018. Expandable group identification in spreadsheets. In *Proceedings of the 33rd ACM/IEEE International Conference on Automated Software Engineering*. 498–508.

[19] Julian Eberius, Katrin Braunschweig, and Others. 2015. Building the dresden web table corpus: A classification approach. In *2015 IEEE/ACM 2nd International Symposium on Big Data Computing (BDC)*. IEEE, 41–50.

[20] Jing Fang, Prasenjit Mitra, Zhi Tang, and C Lee Giles. 2012. Table header detection and classification. In *Twenty-Sixth AAAI Conference on Artificial Intelligence*.

[21] Xi Fang, Weijie Xu, Fiona Anting Tan, Jiani Zhang, Ziqing Hu, Yanjun Qi, Scott Nickleach, Diego Socolinsky, Srinivasan Sengamedu, and Christos Faloutsos. 2024. Large Language Models on Tabular Data—A Survey. *arXiv preprint arXiv:2402.17944* (2024).

[22] Besnik Fetahu, Avishek Anand, and Maria Koutraki. 2019. TableNet: An approach for determining fine-grained relations for wikipedia tables. In *The World Wide Web Conference*. 2736–2742.

[23] Marc Fisher and Gregg Rothmel. 2005. The EUSES spreadsheet corpus: a shared resource for supporting experimentation with spreadsheet dependability mechanisms. In *Proceedings of the first workshop on End-user software engineering*. 1–5.

- [24] Majid Ghasemi Gol, Jay Pujara, and Pedro Szekely. 2019. Tabular Cell Classification Using Pre-Trained Cell Embeddings. In *2019 IEEE International Conference on Data Mining (ICDM)*. IEEE, 230–239.
- [25] Huifeng Guo, Ruiming Tang, Yunming Ye, Zhenguo Li, and Xiuqiang He. 2017. DeepFM: A Factorization-Machine based Neural Network for CTR Prediction. *arXiv:1703.04247* [cs.LG]
- [26] Felienne Hermans and Emerson Murphy-Hill. [n. d.]. Enron's Spreadsheets and Related Emails: A Dataset and Analysis. In *2015 IEEE/ACM 37th IEEE International Conference on Software Engineering*.
- [27] Jonathan Herzig, Thomas Mueller, Syrine Krichene, and Julian Eisenschlos. 2021. Open Domain Question Answering over Tables via Dense Retrieval. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*. 512–519.
- [28] Jonathan Herzig, Pawel Krzysztof Nowak, Thomas Müller, Francesco Piccinno, and Julian Martin Eisenschlos. 2020. TaPas: Weakly Supervised Table Parsing via Pre-training. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics, ACL 2020, Online, July 5-10, 2020*, Dan Jurafsky, Joyce Chai, Natalie Schluter, and Joel R. Tetreault (Eds.). Association for Computational Linguistics, 4320–4333. <https://doi.org/10.18653/V1/2020.ACL-MAIN.398>
- [29] Jinhao Jiang, Kun Zhou, et al. 2023. StructGPT: A General Framework for Large Language Model to Reason over Structured Data. In *EMNLP*.
- [30] Kelvin Jiang, Dekun Wu, and Hui Jiang. 2019. FreebaseQA: A New Factoid QA Data Set Matching Trivia-Style Question-Answer Pairs with Freebase.. In *NAACL-HLT (1)*. 318–323.
- [31] Marcin Kardas, Piotr Czapla, Pontus Stenetorp, Sebastian Ruder, Sebastian Riedel, Ross Taylor, and Robert Stojnic. 2020. AxCell: Automatic Extraction of Results from Machine Learning Papers. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. 8580–8594.
- [32] Elvis Koci, Maik Thiele, Josephine Rehak, Oscar Romero, and Wolfgang Lehner. 2019. DECO: A Dataset of Annotated Spreadsheets for Layout and Table Recognition. In *International Conference on Document Analysis and Recognition*. IEEE.
- [33] Neema Kotonya and Francesca Toni. 2020. Explainable automated fact-checking for public health claims. *arXiv preprint arXiv:2010.09926* (2020).
- [34] Oliver Lehmberg, Dominique Ritze, Robert Meusel, and Christian Bizer. 2016. A large public corpus of web tables containing time and context metadata. In *Proceedings of the 25th International Conference Companion on World Wide Web*. 75–76.
- [35] Hongxin Li, Jingran Su, et al. 2023. SheetCopilot: Bringing Software Productivity to the Next Level through Large Language Models. In *NeurIPS*.
- [36] Peng Li, Yeye He, Dror Yashar, Weiwei Cui, Song Ge, Haidong Zhang, Danielle Rifinski Fainman, Dongmei Zhang, and Surajit Chaudhuri. 2023. Tablegpt: Table-tuned gpt for diverse table tasks. *arXiv preprint arXiv:2310.09263* (2023).
- [37] Weizhe Lin, Rexhina Blloshmi, Bill Byrne, et al. 2023. An inner table retriever for robust table question answering. In *ACL*.
- [38] Weizhe Lin, Rexhina Blloshmi, Bill Byrne, et al. 2023. LI-RAGE: Late interaction retrieval augmented generation with explicit signals for open-domain table question answering. In *ACL*.
- [39] Qian Liu, Bei Chen, Jiaqi Guo, Morteza Ziyadi, Zeqi Lin, Weizhu Chen, and Jian-Guang Lou. 2022. TAPEX: Table Pre-training via Learning a Neural SQL Executor. In *The Tenth International Conference on Learning Representations, ICLR 2022*. <https://openreview.net/forum?id=O50443AsCP>
- [40] Pan Lu, Liang Qiu, Kai-Wei Chang, Ying Nian Wu, Song-Chun Zhu, Tanmay Rajpurohit, Peter Clark, and Ashwin Kalyan. 2023. Dynamic Prompt Learning via Policy Gradient for Semi-structured Mathematical Reasoning. *arXiv:2209.14610* [cs.LG]
- [41] Weizheng Lu, Jiaming Zhang, Jing Zhang, and Yueguo Chen. 2024. Large Language Model for Table Processing: A Survey. *arXiv preprint arXiv:2402.05121* (2024).
- [42] Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegrefe, Uri Alon, Nouha Dziri, Shrimai Prabhumoye, Yiming Yang, et al. 2024. Self-refine: Iterative refinement with self-feedback. *Advances in Neural Information Processing Systems* 36 (2024).
- [43] Ankur Parikh, Xuezhi Wang, et al. 2020. ToTTo: A Controlled Table-To-Text Generation Dataset. In *EMNLP*.
- [44] Panupong Pasupat and Percy Liang. 2015. Compositional Semantic Parsing on Semi-Structured Tables. In *Proc. of ACL*.
- [45] Panupong Pasupat and Percy Liang. 2015. Compositional Semantic Parsing on Semi-Structured Tables. *arXiv:1508.00305* [cs.CL]
- [46] Mohammadreza Pourreza and Davood Rafiei. 2023. DIN-SQL: Decomposed In-Context Learning of Text-to-SQL with Self-Correction. *arXiv:2304.11015* [cs.CL]
- [47] Jay Pujara, Pedro Szekely, Huan Sun, and Muhao Chen. 2021. From tables to knowledge: Recent advances in table understanding. In *Proceedings of the 27th ACM SIGKDD Conference on Knowledge Discovery & Data Mining*. 4060–4061.
- [48] Qi Shi, Yu Zhang, Qingyu Yin, and Ting Liu. 2021. Logic-level Evidence Retrieval and Graph-based Verification Network for Table-based Fact Verification. *arXiv preprint arXiv:2109.06480* (2021).
- [49] Alexey Shigarov, Andrey Mikhailov, and Andrey Altaev. 2016. Configurable table structure recognition in untagged pdf documents. In *Proceedings of the 2016 ACM Symposium on Document Engineering*. 119–122.
- [50] Ananya Singha, José Cambronero, et al. 2023. Tabular Representation, Noisy Operators, and Impacts on Table Structure Understanding Tasks in LLMs. In *NeurIPS TRL Workshop*.
- [51] Yuan Sui, Mengyu Zhou, Mingjie Zhou, Shi Han, and Dongmei Zhang. 2023. GPT4Table: Can Large Language Models Understand Structured Table Data? A Benchmark and Empirical Study. *arXiv preprint arXiv:2305.13062*.
- [52] Yuan Sui, Jiaru Zou, Mengyu Zhou, Xinyi He, Lun Du, Shi Han, and Dongmei Zhang. 2023. TAP4LLM: Table Provider on Sampling, Augmenting, and Packing Semi-structured Data for Large Language Model Reasoning. *arXiv:2312.09039* [cs.CL]
- [53] Tianshu Wang, Hongyu Lin, Xianpei Han, Le Sun, Xiaoyang Chen, Hao Wang, and Zhenyu Zeng. 2023. DBCopilot: Scaling Natural Language Querying to Massive Databases. *arXiv:2312.03463* [cs.CL]
- [54] Xuezhi Wang, Jason Wei, et al. 2023. Self-Consistency Improves Chain of Thought Reasoning in Language Models. In *ICLR*.
- [55] Zhiruo Wang, Daniel Fried, and Graham Neubig. 2024. TroVE: Inducing Verifiable and Efficient Toolboxes for Solving Programmatic Tasks. *arXiv preprint arXiv:2401.12869* (2024).
- [56] Zhiruo Wang, Zhengbao Jiang, Eric Nyberg, and Graham Neubig. 2022. Table Retrieval May Not Necessitate Table-specific Model Design. In *Proceedings of the Workshop on Structured and Unstructured Knowledge Integration (SUKI)*. Association for Computational Linguistics.
- [57] Zilong Wang, Hao Zhang, Chun-Liang Li, Julian Martin Eisenschlos, Vincent Perot, Zifeng Wang, Lesly Miculicich, Yasuhisa Fujii, Jingbo Shang, Chen-Yu Lee, et al. 2024. Chain-of-table: Evolving tables in the reasoning chain for table understanding. *arXiv preprint arXiv:2401.04398* (2024).
- [58] Jason Wei, Maarten Bosma, Vincent Zhao, Kelvin Guu, Adams Wei Yu, Brian Lester, Nan Du, Andrew M. Dai, and Quoc V Le. 2022. Finetuned Language Models are Zero-Shot Learners. In *International Conference on Learning Representations*. <https://openreview.net/forum?id=gEzRGCozdqR>
- [59] Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, Denny Zhou, et al. 2022. Chain-of-thought prompting elicits reasoning in large language models. *Advances in Neural Information Processing Systems* 35 (2022), 24824–24837.
- [60] Siqiao Xue, Caigao Jiang, Wenhui Shi, et al. 2023. DB-GPT: Empowering Database Interactions with Private Large Language Models. *arXiv preprint arXiv:2312.17449* (2023).
- [61] Shunyu Yao, Jeffrey Zhao, et al. 2023. ReAct: Synergizing Reasoning and Acting in Language Models. In *ICLR*.
- [62] Yunhu Ye, Binyuan Hui, et al. 2023. Large Language Models Are Versatile Decomposers: Decomposing Evidence and Questions for Table-based Reasoning. In *SIGIR*.
- [63] Tao Yu, Rui Zhang, Kai Yang, Michihiro Yasunaga, Dongxu Wang, Zifan Li, James Ma, Irene Li, Qingning Yao, Shanelle Roman, et al. 2018. Spider: A large-scale human-labeled dataset for complex and cross-domain semantic parsing and text-to-sql task. *arXiv preprint arXiv:1809.08887* (2018).
- [64] Lifan Yuan, Yangyi Chen, Xingyao Wang, Yi Fung, Hao Peng, and Heng Ji. 2024. CRAFT: Customizing LLMs by Creating and Retrieving from Specialized Toolsets. In *The Twelfth International Conference on Learning Representations*. <https://openreview.net/forum?id=GövdSt9XM>
- [65] Hongzhi Zhang, Yingyao Wang, Sirui Wang, Xuezhi Cao, Fuzheng Zhang, and Zhongyuan Wang. 2020. Table fact verification with structure-aware transformer. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. 1624–1629.
- [66] Li Zhang, Shuo Zhang, and Krisztian Balog. 2019. Table2vec: Neural word and entity embeddings for table population and retrieval. In *Proceedings of International ACM SIGIR Conference on Research and Development in Information Retrieval*.
- [67] Shuo Zhang and Krisztian Balog. 2020. Web Table Extraction, Retrieval, and Augmentation: A Survey. *ACM Trans. Intell. Syst. Technol.* 11, 2, Article 13 (Jan. 2020), 35 pages. <https://doi.org/10.1145/3372117>
- [68] Tianshu Zhang, Xiang Yue, Yifei Li, and Huan Sun. 2023. TableLLama: Towards Open Large Generalist Models for Tables. *arXiv:2311.09206* [cs.CL]
- [69] Wenqi Zhang, Yongliang Shen, et al. 2023. Data-Copilot: Bridging Billions of Data and Humans with Autonomous Workflow. *arXiv preprint arXiv:2306.07209* (2023).
- [70] Xuanliang Zhang, Dingzirui Wang, Longxu Dou, Qingfu Zhu, and Wanxiang Che. 2024. A Survey of Table Reasoning with Large Language Models. *arXiv preprint arXiv:2402.08259* (2024).
- [71] Yunjia Zhang, Jordan Henkel, Avriella Floratou, Joyce Cahoon, Shaleen Deep, and Jignesh M. Patel. 2023. ReActTable: Enhancing ReAct for Table Question Answering. *arXiv:2310.00815* [cs.DB]
- [72] Guangzhen Zhao and Peng Yang. 2022. Table-based Fact Verification with Self-labeled Keypoint Alignment. In *Proc. of COLING*.