

11-667 Large Language Models: Methods and Applications

Retrieval and Retrieval-Augmented Generation

Akari Asai


aasai@andrew.cmu.edu | akaria@allenai.org
<https://akariasai.github.io/>



Slides adapted from

ACL 2023 tutorial by Akari Asai, Sewon Min, Zexuan Zhong, Danqi Chen <https://acl2023-retrieval-lm.github.io/>
Advanced NLP Fall 2024 by Prof. Graham Neubig <https://phontron.com/class/anlp-fall2024/>

Parametric LMs



Limitations of Current LMs: Hallucinations

Explain Akari Asai's research, including her representative papers

(ChatGPT4o, February 8, 2025)

Akari Asai is a researcher in natural language processing (NLP) and artificial intelligence (AI), particularly focusing on open-domain question answering, knowledge retrieval. Her work primarily aims to improve how machines retrieve, understand, and r information efficiently.

Learning to Retrieve Passages without Supervision

Ori Ram Gal Shachaf Omer Levy Jonathan Berant Amir Globerson

Blavatnik School of Computer Science, Tel Aviv University
ori.ram@cs.tau.ac.il

Representative Papers:

NAACL 2022

1. "Learning to Retrieve Passages without Supervision" (~~NeurIPS 2022~~)

- Introduces a method for training retrieval models without labeled data, improving open-domain QA efficiency.
- Leverages contrastive learning to align questions and relevant passages.

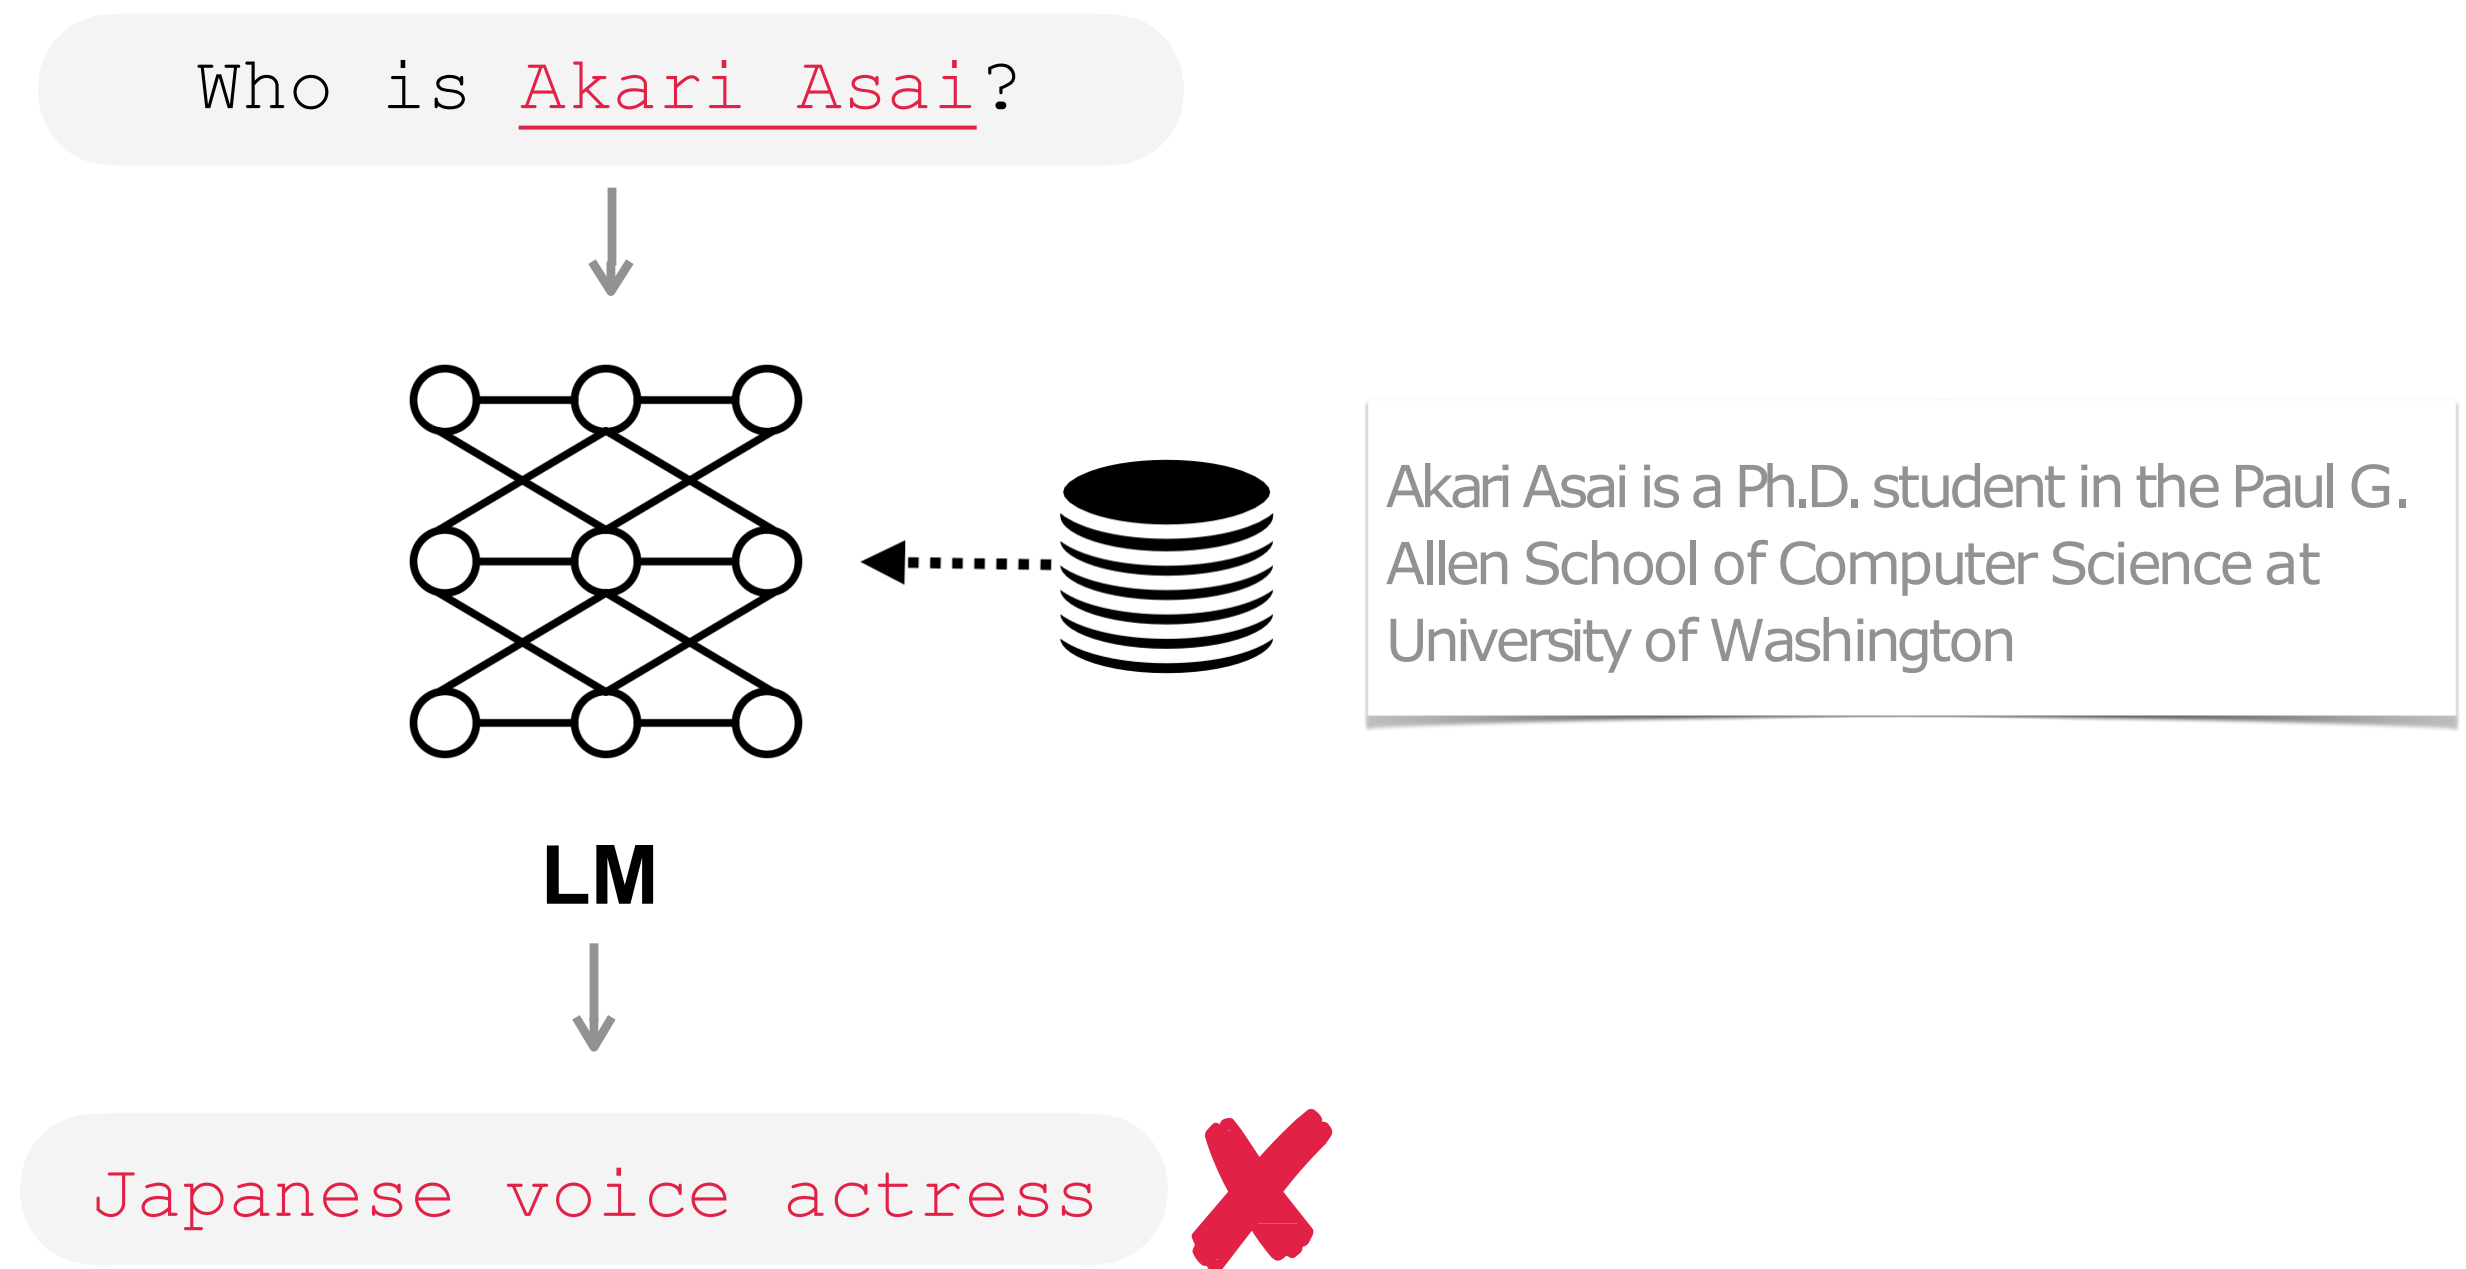
 **Not my paper**

 **Venue is wrong**

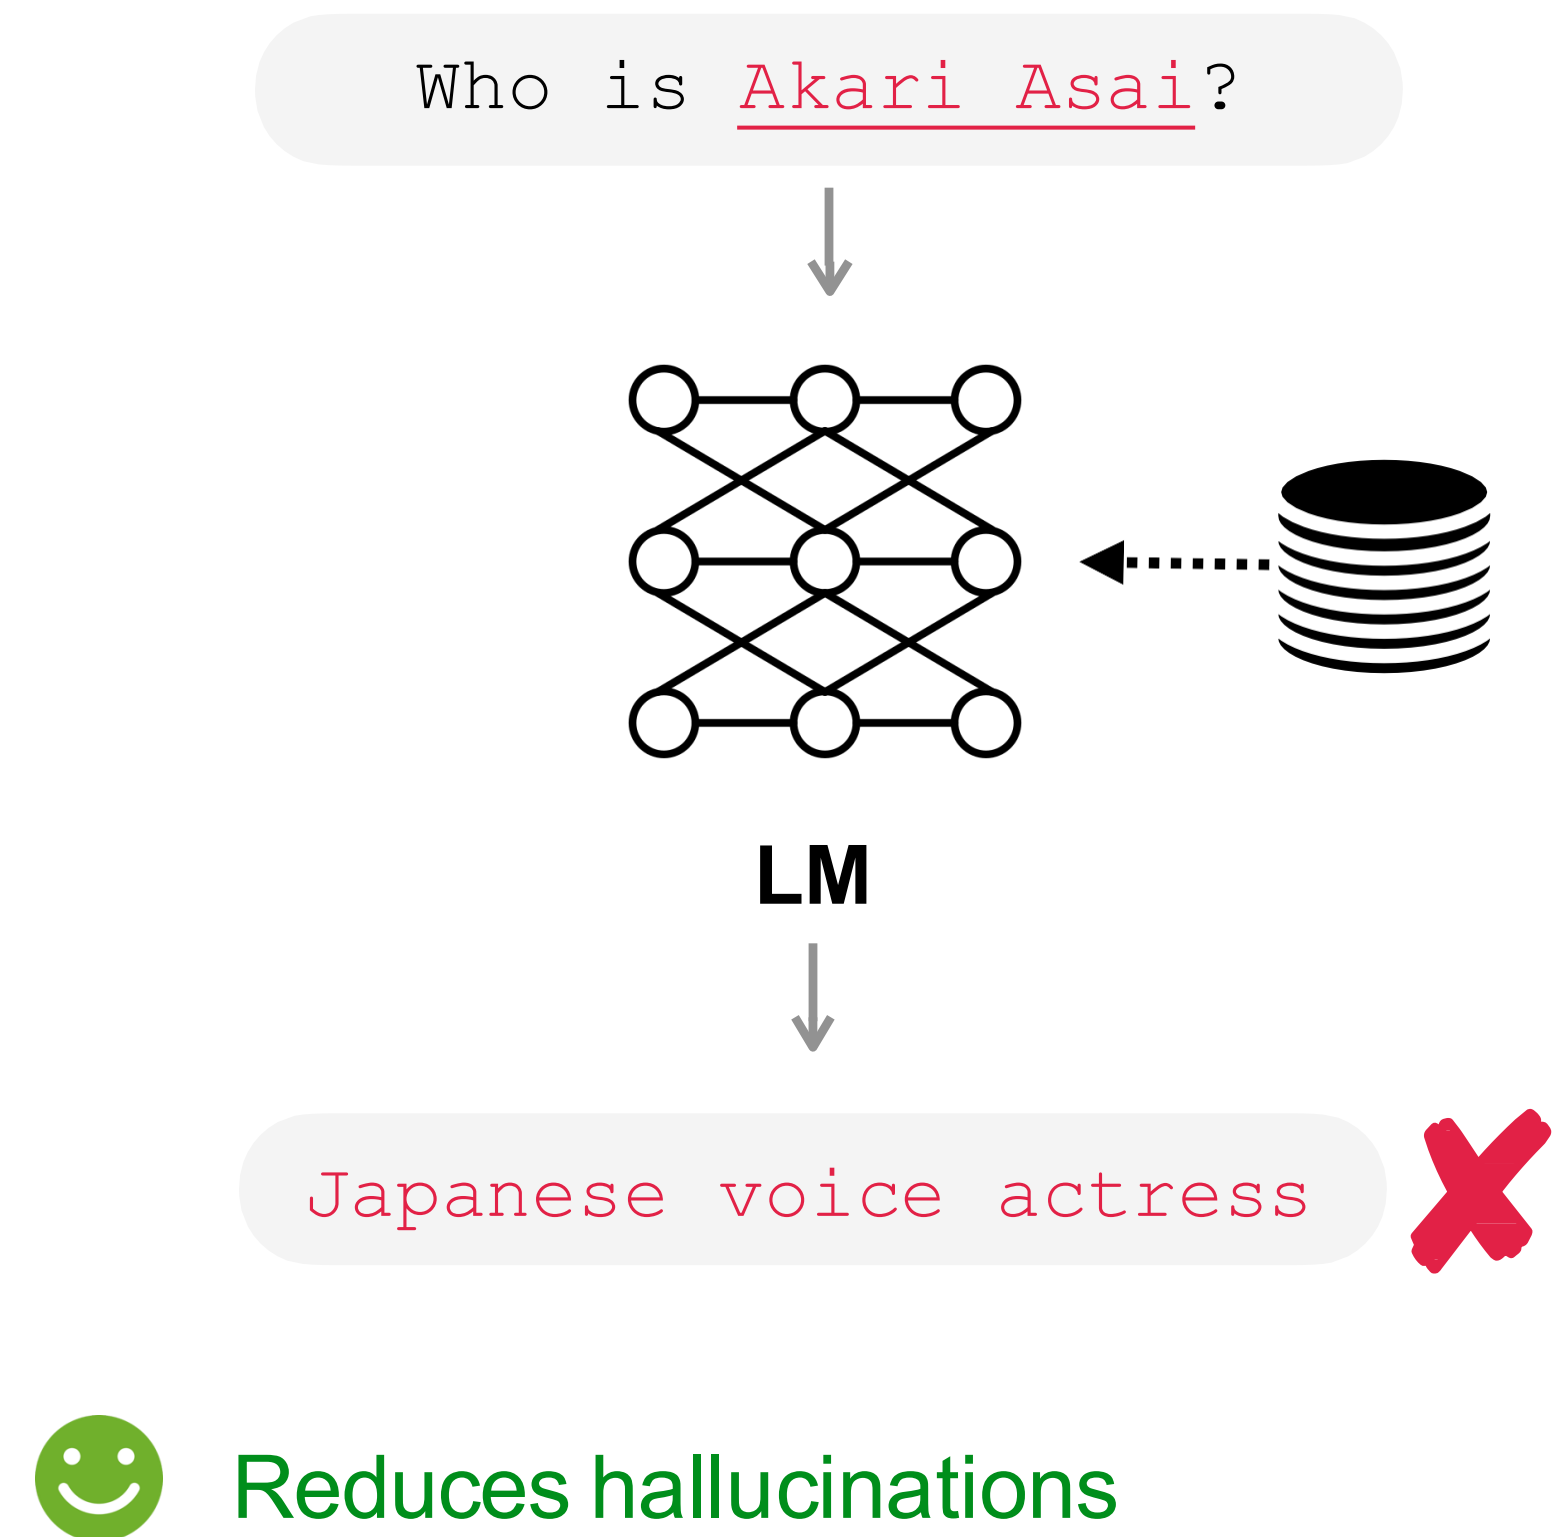

2. ~~"One Question, Many Answers: A Retrieval-based Multimodal QA Dataset"~~ (EMNLP 2022)

LMs struggle in long-tail knowledge

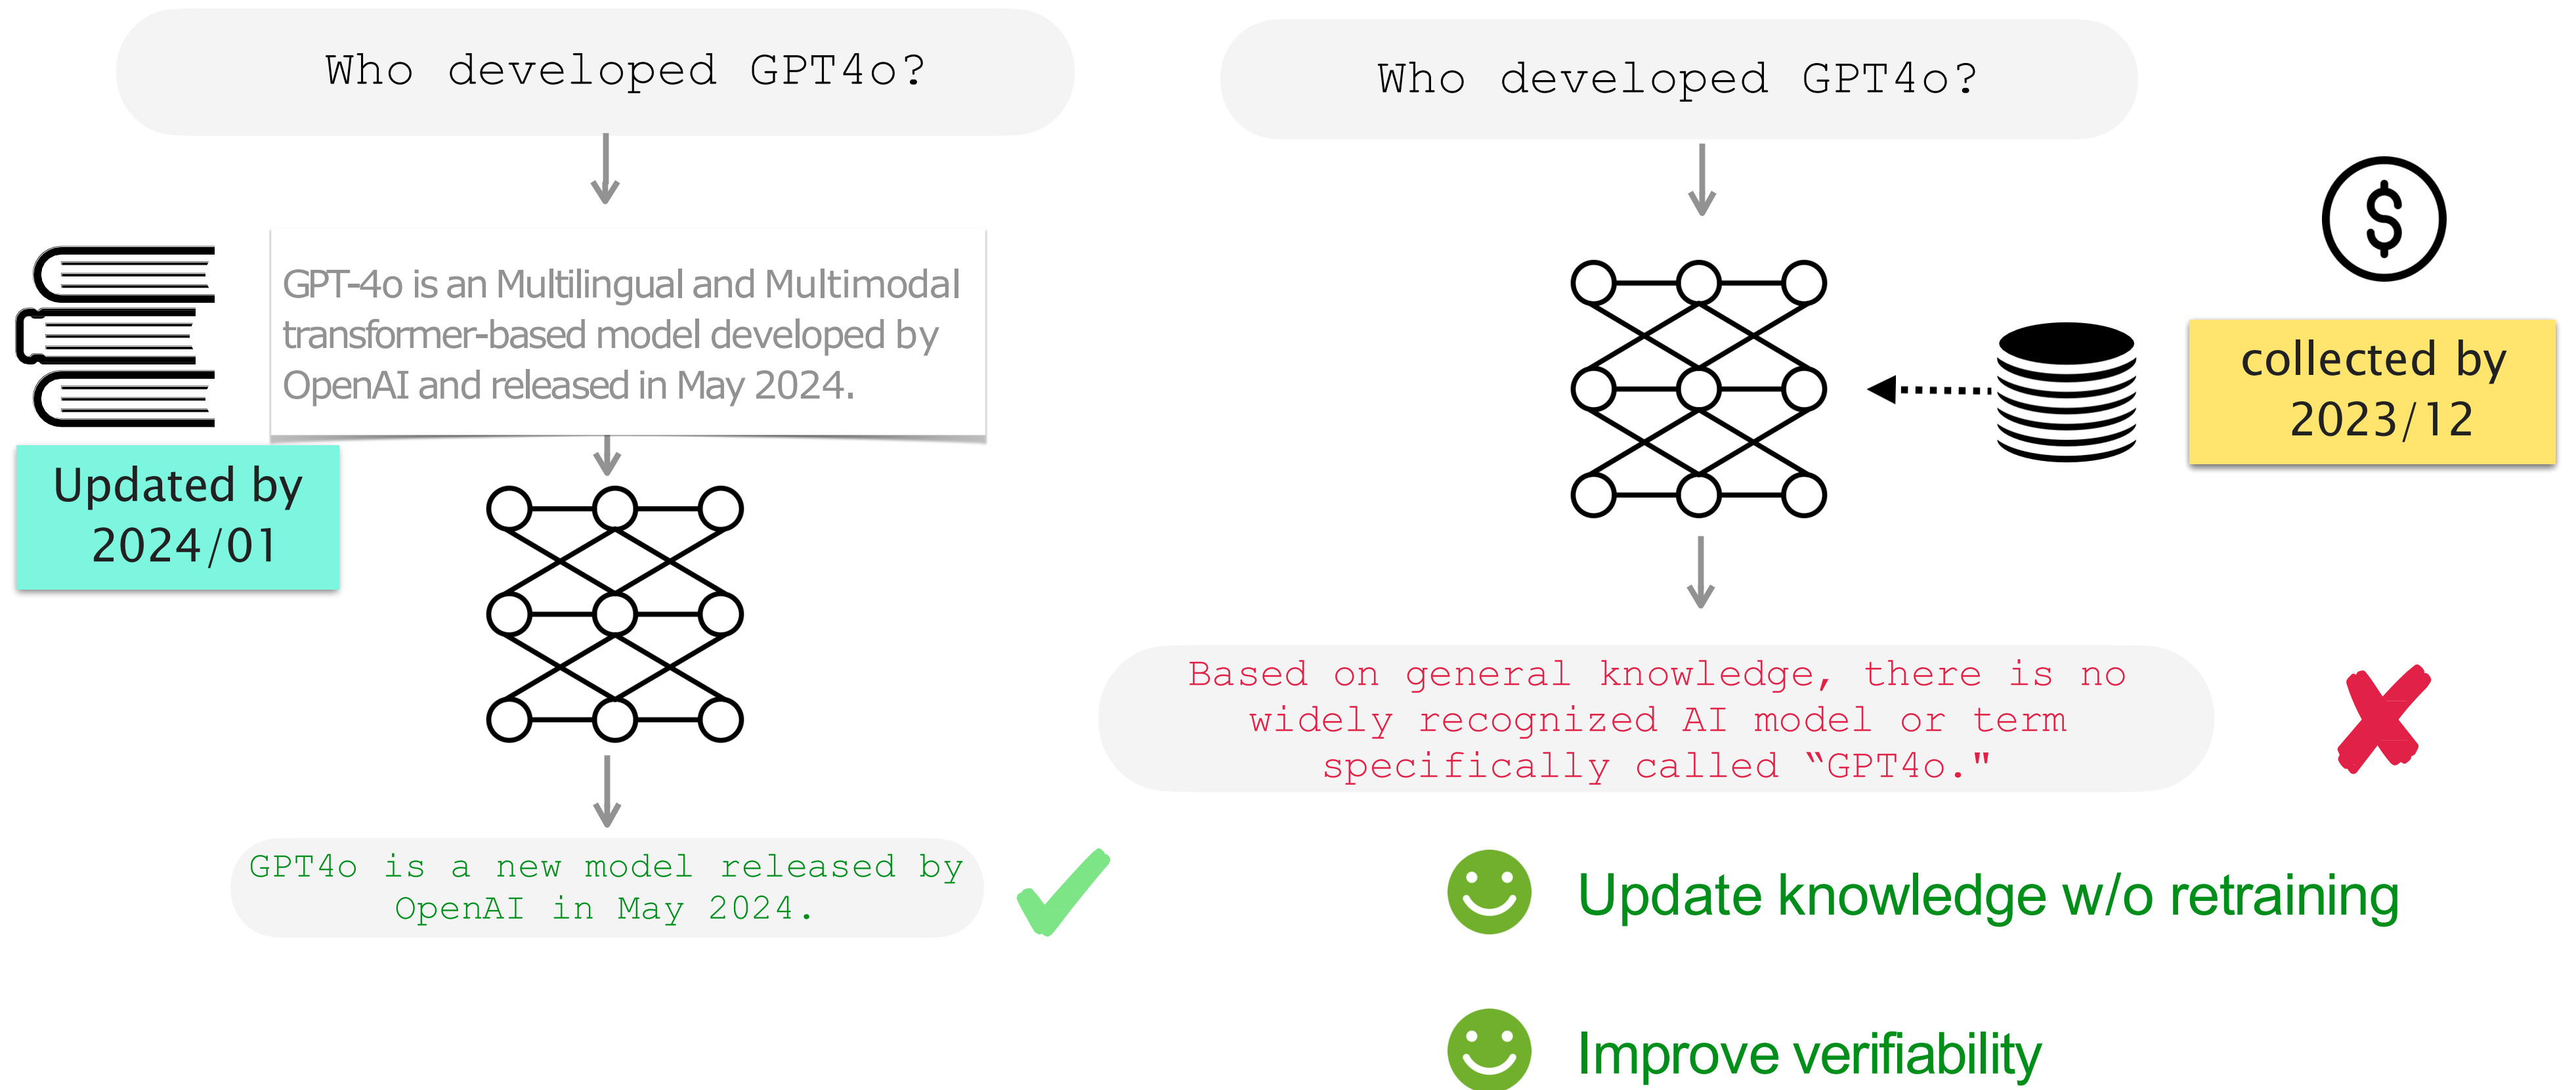
Retrieval-Augmented LMs: Intuition




Retrieval-Augmented LMs: Intuition



Retrieval-Augmented LMs: Intuition




Widespread Adoptions in Real World



At Databricks 60% of LLM applications use some form of **retrieval-augmented generation (RAG)**

The Shift from Models to Compound AI Systems

Retrieval-Augmented LMs: Overview




Datastore

Collections of a large number of documents



Retriever

Retrieve top k documents in datastore




LM

GPT-4o is a pre-trained transformer developed by OpenAI.

Transformers is a series of science fiction action films based on the Transformers franchise.

GPT4o was released by OpenAI in May 2024.

Retrieval-Augmented LMs: Overview



$\text{Sim}(\cdot | x)$

GPT-4o is a pre-trained transformer developed by OpenAI.

0.9


Transformers is a series of science fiction action films based on the Transformers franchise.

0.1

GPT4o was released by OpenAI in May 2024.

0.8

Retrieval-Augmented LMs: Overview



$$D \in \text{Top}_k \text{Sim}(\cdot | x)$$

GPT-4o is a pre-trained transformer developed by OpenAI.

0.9


Transformers is a series of science fiction action films based on the Transformers franchise.

0.1

GPT4o was released by OpenAI in May 2024.

0.8

Retrieval-Augmented LMs: Overview



$$D \in \text{Top}_k \text{Sim}(\cdot | x)$$

Transformers is a series of science fiction action films based on the Transformers franchise.

GPT-4o is a pre-trained transformer developed by OpenAI.


0.9

0.1

GPT4o was released by OpenAI in May 2024.


0.8

Today's Outline




- ✓ Sources of datastore
- ✓ Processing
- ✓ Scaling

Today's Outline




- ✓ Types of retrievers
- ✓ Training
- ✓ Evaluations

Today's Outline



- ✓ Common architectures
- ✓ Recent progress in RAG

Today's Outline



- ✓ Sources of datastore
- ✓ Processing
- ✓ Scaling

What Should Be Used as “Datastore”?

Which company developed GPT4o?

English Wikipedia




Chen et al., 2017; Gu et al., 2020; Asai et al., 2020;
Guu et al., 2021; Lewis et al., 2021 ... etc

<https://dumps.wikimedia.org/>

How should I implement RAG
using LlamaIndex?


Code snippets



Official documentations



LangChain



Community forums




Massively Scaling Datastore

MassiveDS

1.4 trillion tokens (22TB)



Datastore Scaling



Processing Documents

☰ GPT-4	🌐 32 languages ▼
Article Talk	Read Edit View history Tools ▼
From Wikipedia, the free encyclopedia	
Generative Pre-trained Transformer 4 (GPT-4) is a multimodal large language model trained and created by OpenAI and the fourth in its series of GPT foundation models.^[1] It was launched on March 14, 2023,^[1] and made publicly available via the paid chatbot product ChatGPT Plus, via OpenAI's API, and via the free chatbot Microsoft Copilot.^[2] As a transformer-based model, GPT-4 uses a paradigm where pre-training using both public data and "data licensed from third-party providers" is used to predict the next token. After this step, the model was then fine-tuned with reinforcement learning feedback from humans and AI for human alignment and policy compliance.^{[3]:2} Observers reported that the iteration of ChatGPT using GPT-4 was an improvement on the previous iteration based on GPT-3.5, with the caveat that GPT-4 retains some of the problems with earlier revisions.^[4] GPT-4, equipped with vision capabilities (GPT-4V),^[5] is capable of taking images as input on ChatGPT.^[6] OpenAI has not revealed technical details and statistics about GPT-4, such as the precise size of the model.^[7]	Generative Pre-trained Transformer 4 (GPT-4) Developer(s) OpenAI Initial release March 14, 2023; 22 months ago Predecessor GPT-3.5 Successor GPT-4o Type Multimodal Large language model Generative pre-trained transformer Foundation model License Proprietary Website openai.com/gpt-4 ↗
Background [edit]	
<i>Further information:</i> GPT-3 § Background, and GPT-2 § Background	
OpenAI introduced the first GPT model (GPT-1) in 2018, publishing a paper called "Improving Language Understanding by Generative Pre-	Part of a series on Machine learning

<https://en.wikipedia.org/wiki/GPT-4>

Processing Documents

Curate and preprocess data

e.g., HTML -> Plain text



Chunking

Paragraph-level (e.g., \n)

Every k words (e.g., 100-250)

GPT-4o is a pre-trained transformer developed by OpenAI.

Transformers is a series of science fiction action films based on the Transformers franchise.

GPT4o was released by OpenAI in May 2024.

@I\$O@

Post-processing


e.g., Remove short documents

GPT-4o is a pre-trained transformer developed by OpenAI.

Transformers is a series of science fiction action films based on the Transformers franchise.

GPT4o was released by OpenAI in May 2024.


Summary of Part 1



- ✓ Sources of datastore
- ✓ Processing
- ✓ Scaling

- Choosing **the right datastore** is important
- **Chunking** and **filtering** strategies are important
- **Scaling** datastores offer performance gain while adding technical challenges

Today's Outline



- ✓ Types of retrievers
- ✓ Training
- ✓ Evaluations

Types of Retrievers

$$D \in \text{Top}_k \text{Sim}(\cdot | x)$$

Sparse retrievers

- **Sim:** Term-frequency based embeddings
- Training is not required


e.g., TF-IDF, BM25



Dense retrievers

- **Sim:** dense embeddings encoded by pre-trained LMs
- Training is needed*

e.g., DPR, Contriever, ColBERT



Rerankers

- **Sim:** Scores based on jointly encoded query and doc
- Training is needed*

e.g., cross-encoder reranker



Sparse Retrievers: One-hot Vectors

	q=what is nlp	d ₁ = what is life ? candy is life !	d ₂ = nlp is an acronym for natural language processing	d ₃ = I like to do good research on nlp
what	1	1	0	0
candy	0	1	0	0
nlp	1	0	1	1
is	1	1	1	0
language	0	0	0	0
life	0	1	0	0
...	...	...	...	...

Check if a term appears in a document

Sparse Retrievers: Term-count Vectors

	q=what is nlp	d ₁ = what is life ? candy is life !	d ₂ = nlp is an acronym for natural language processing	d ₃ = I like to do good research on nlp
what	1	1	0	0
candy	0	1	0	0
nlp	1	0	1	1
is	1	1	1	0
language	0	0	0	0
life	0	2	0	0
...	...	...	...	...

Count the number of appearances in a doc

Sparse Retrievers: Computing Weighted Term Scores

$d_1 = \text{what is life ?}$
 candy is life !

$\text{TF}(t, d) = \frac{\text{freq}(t, d)}{\sum_{t'} \text{freq}(t', d)}$

$\uparrow$
 $t_1 = \text{what}$

of documents

$\text{IDF}(t) = \log \frac{|D|}{\sum_{d' \in D} \delta(\text{freq}(t, d') > 0)}$

of documents where term t appears

$$\text{TF-IDF}(t, d) = \text{TF}(t, d) \times \text{IDF}(t)$$

$$\text{BM-25}(t, d) = \text{IDF}(t) \cdot \frac{\text{freq}(t, d) \cdot (k_1 + 1)}{\text{freq}(t, d) + k_1 \cdot \left(1 - b + b \cdot \frac{|d|}{\text{avgdl}}\right)}$$

Sparse Retrievers: Weighted-term Vectors

	q=what is nlp	d ₁ = what is life ? candy is life !	d ₂ = nlp is an acronym for natural language processing	d ₃ = I like to do good research on nlp
what	0.36	0.18	0	0
candy	0	0.18	0	0
nlp	0.13	0	0.05	0.05
is	0.13	0.13	0.05	0
language	0	0	0.13	0
life	0	0.36	0	0
...	...	...	...	...

Compute TF-IDF weights to build weighted vectors

Sparse Retrievers: Weighted-term Vectors

	q=what is nlp	d ₁ = what is life ? candy is life !	d ₂ = nlp is an acronym for natural language processing	d ₃ = I like to do good research on nlp
what	0.36	0.18	0	0
candy	0	0.18	0	0
nlp	0.13	0	0.05	0.05
is	0.13	0.13	0.05	0
language	0	0	0.13	0
life	0	0.36	0	0
...	...	...	...	...

Compute TF-IDF weights to build weighted vectors

Sparse Retrievers: Weighted-term Vectors

q=what is nlp

what	0.36
candy	0
nlp	0.13
is	0.13
language	0
life	0
...	...

d_1 = what is life ?
candy is life !

0.18
0.18
0
0.13
0
0.36
...

d_2 = nlp is an
acronym for natural
language processing

0
0
0.05
0.05
0.13
0
...

d_3 = I like to do
good research
on nlp

0
0
0.05
0
0
0
...

Compute cosine similarity

$$q * d_1 = 0.44$$

$$q * d_2 = 0.21$$

$$q * d_3 = 0.32$$


Computing TF-IDF Matrices: Weighted-term Vectors

q=what is nlp	$d_1 = \text{what is life ?}$ candy is life !	$d_2 = \text{nlp is an}$ $\text{acronym for natural}$ $\text{language processing}$	$d_3 = \text{I like to do}$ good research on nlp
what	0.18	0	0
candy	0.18	0	0
nlp	0	0.05	0.05
is	0.13	0.05	0
language	0	0.13	0
life	0.36	0	0
...	...	...	...
“Bag-of-words”	$q * d_1 = 0.44$	$q * d_2 = 0.21$	$q * d_3 = 0.32$

Can't fully capture semantic similarities

Dense Retrievers: Overview

$\mathbf{x}$ = How many teams will participate in FIFA World Cup 2026?



$$z_1, \dots, z_k = \text{argTop-}k(\mathbf{x} \cdot \mathbf{z})$$


k retrieved chunks

Dense Retrievers: Generating *Embeddings*

- Use output of [CLS] token in masked LMs

e.g., DPR

$\mathbb{R}^d$



- Mean / Max pooling of output vectors

e.g., SBERT, SGPT

$\mathbb{R}^{N \times d}$

	NLI	STSb
<i>Pooling Strategy</i>		
MEAN	80.78	87.44
MAX	79.07	69.92
CLS	79.80	86.62

Karpukhin et al. EMNLP 2020. Dense Passage Retrieval for Open-Domain Question Answering.
Reimers et al. EMNLP 2019. Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks.

Fast Nearest Neighbor Search

Method	Class name	index_factory	Main parameters	Bytes/vector	Exhausti
Exact Search for L2	IndexFlatL2	"Flat"	d	4*d	yes
Exact Search for Inner Product	IndexFlatIP	"Flat"	d	4*d	yes
Hierarchical Navigable Small World graph exploration	IndexHNSWFlat	"HNSW,Flat"	d, M	4*d + x * M * 2 * 4	no
Inverted file with exact post-verification	IndexIVFFlat	"IVFx,Flat"	quantizer, d, nlists, metric	4*d + 8	no
Locality-Sensitive Hashing (binary flat index)	IndexLSH	-	d, nbits	ceil(nbbits/8)	yes
Scalar quantizer (SQ) in flat mode	IndexScalarQuantizer	"SQ8"	d	d	yes
Product quantizer (PQ) in flat mode	IndexPQ	"PQx", "PQ"M"x"nbits	d, M, nbits	ceil(M * nbits / 8)	yes
IVF and scalar quantizer	IndexIVFScalarQuantizer	"IVFx,SQ4" "IVFx,SQ8"	quantizer, d, nlists, qtype	SQfp16: 2 * d + 8, SQ8: d + 8 or SQ4: d/2 + 8	no
IVFADC (coarse quantizer+PQ on residuals)	IndexIVFPQ	"IVFx,PQ"y"x"nbits	quantizer, d, nlists, M, nbits	ceil(M * nbits/8)+8	no

Exact search (still fast for $10^6 \sim 10^7$ scale)


Approximate search (faster but more memory)

Reduce index size with quantization

<https://github.com/facebookresearch/faiss/wiki>

https://speakerdeck.com/matsui_528/cvpr20-tutorial-billion-scale-approximate-nearest-neighbor-search
(CVPR 2020 Tutorial)

Training Dense Retriever

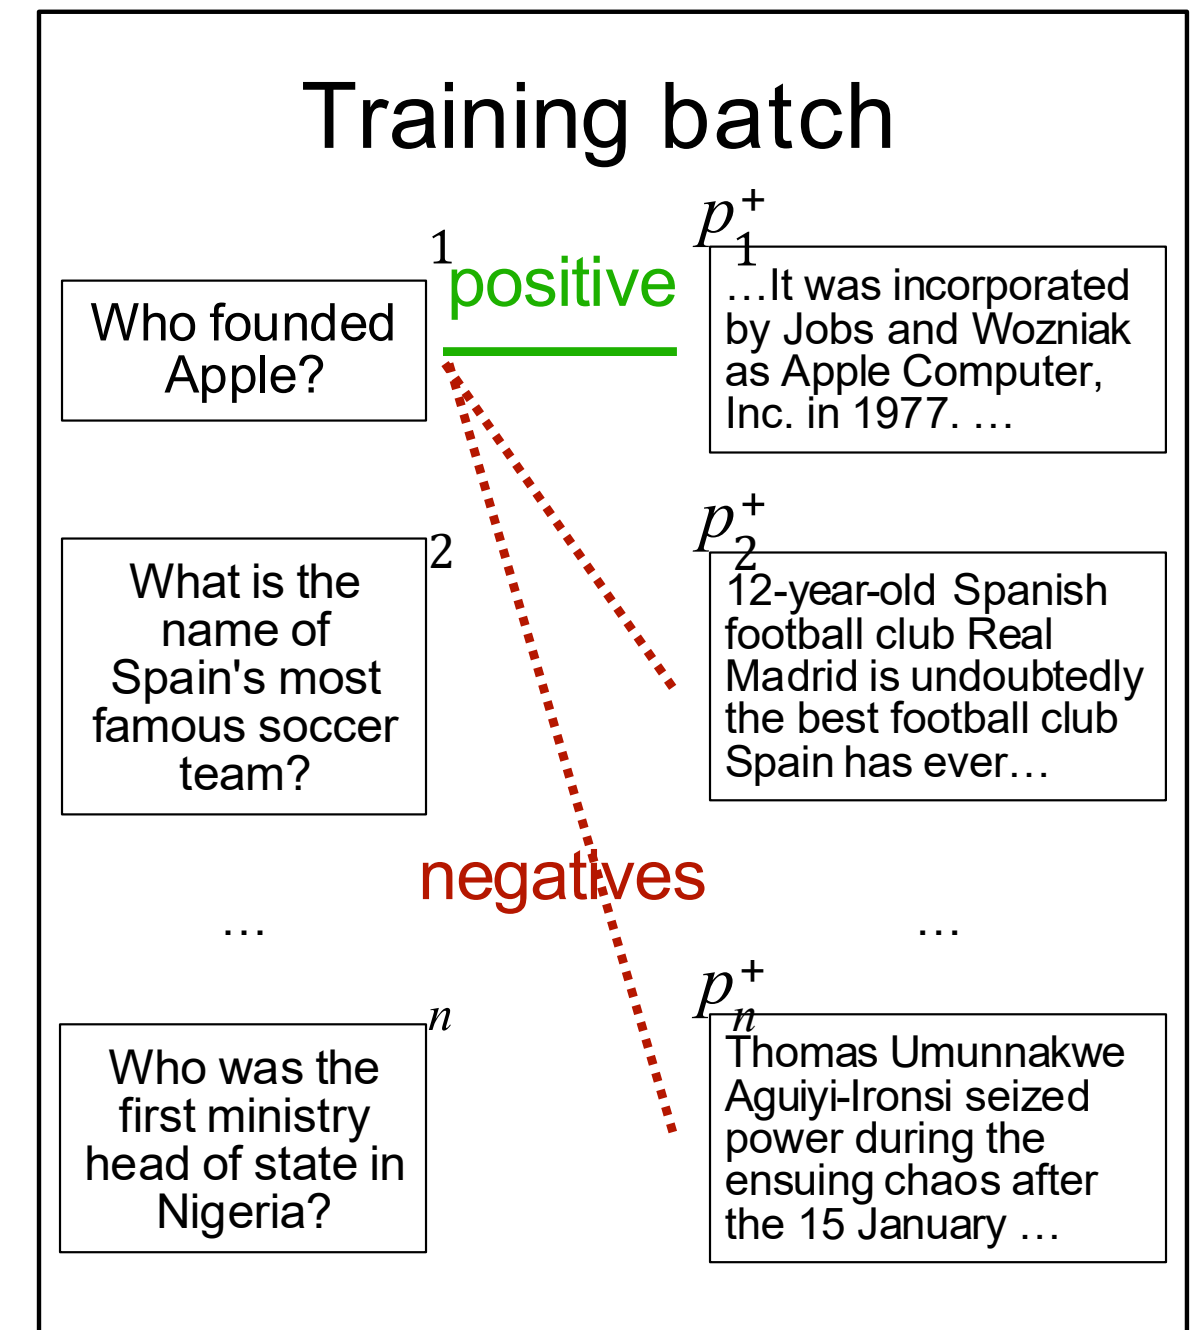


Training Dense Retriever

$$L(q, p^+, p_1^-, p_2^-, \dots, p_n^-)$$
$$= -\log \frac{\exp(\text{sim}(q, p^+))}{\exp(\text{sim}(q, p^+)) + \sum_{j=1}^n \exp(\text{sim}(q, p_j^-))}$$


In-batch negatives

Hard negative retrieved by the same / another model



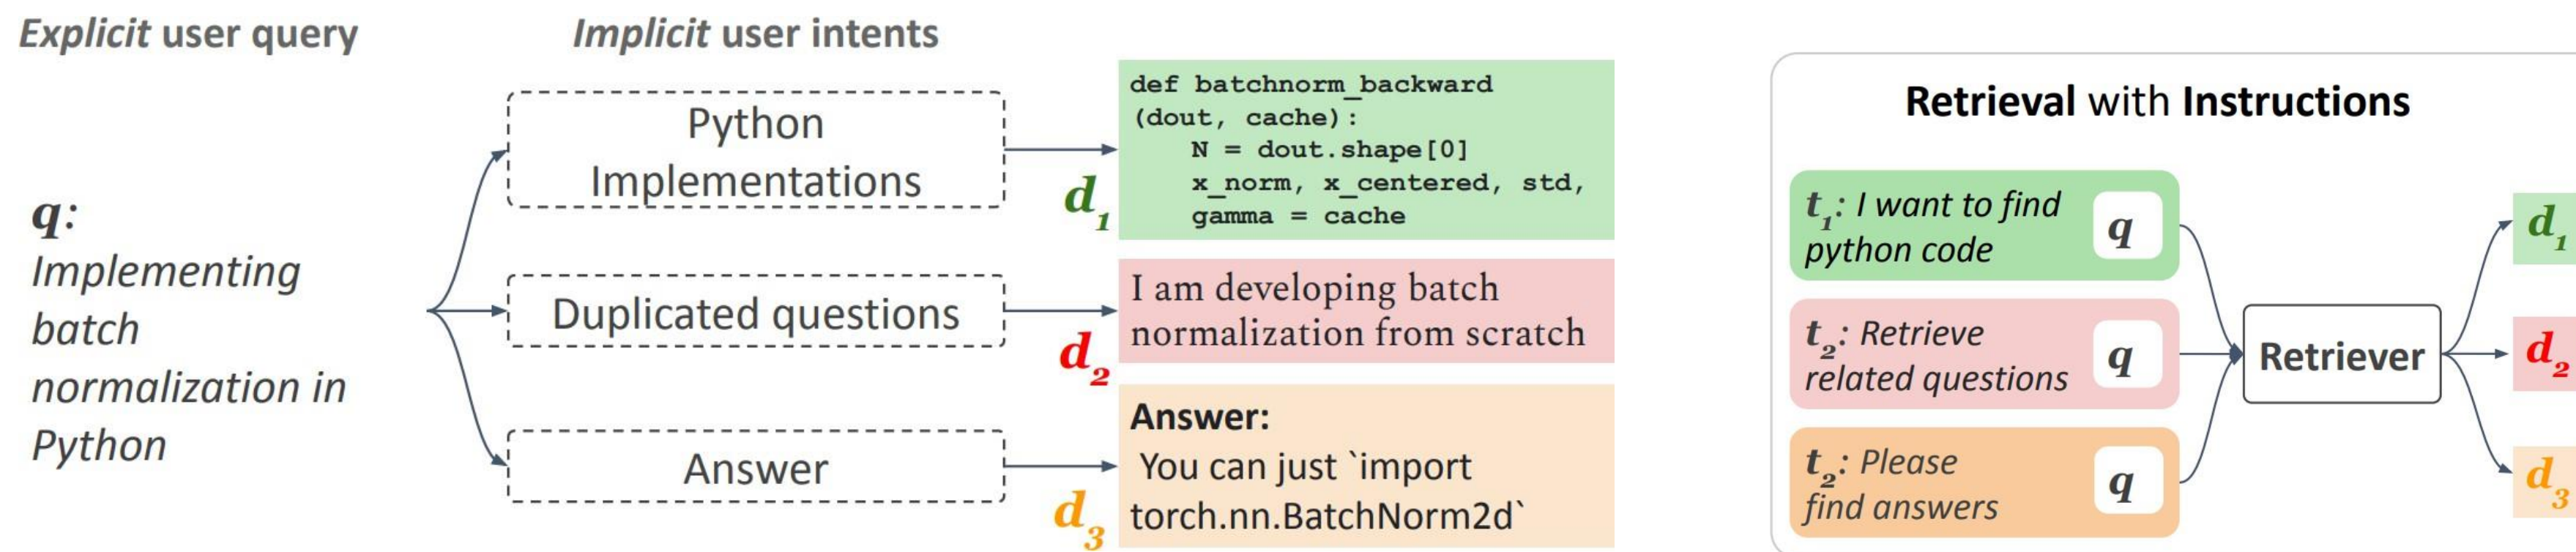
Unsupervised Training

Independent Cropping



Unsupervised dense retrieval model!

Instruction Tuning for Retriever



Dup. Question Retrieval

t_1 : Retrieve a question asked in StackOverflow similar to this

q: How to compute square root in iOS?

How can we calculate the square root in Objective C or Swift?
StackOverflow Question

Dialogue Response Retrieval

t_1 : Find an informative dialogue response to this user's conversation

q: Are armadillos native to a Spanish-speaking part of the world?

Yes, they are most commonly found in North, Central, and South America.
Dialogue Response

Hard negative documents d^{HD}

Which python function can I use to compute sq root?
StackOverflow Question

I love animals and think armadillos are awesome with their leathery shell.
Dialogue Response

Instruction-unfollowing negatives d^{UF}

You can just use the Objective C or Swift's **sqrt** function
StackOverflow Answer

Armadillos are medium-sized mammals found in North, Central, and South America
Wikipedia Paragraph

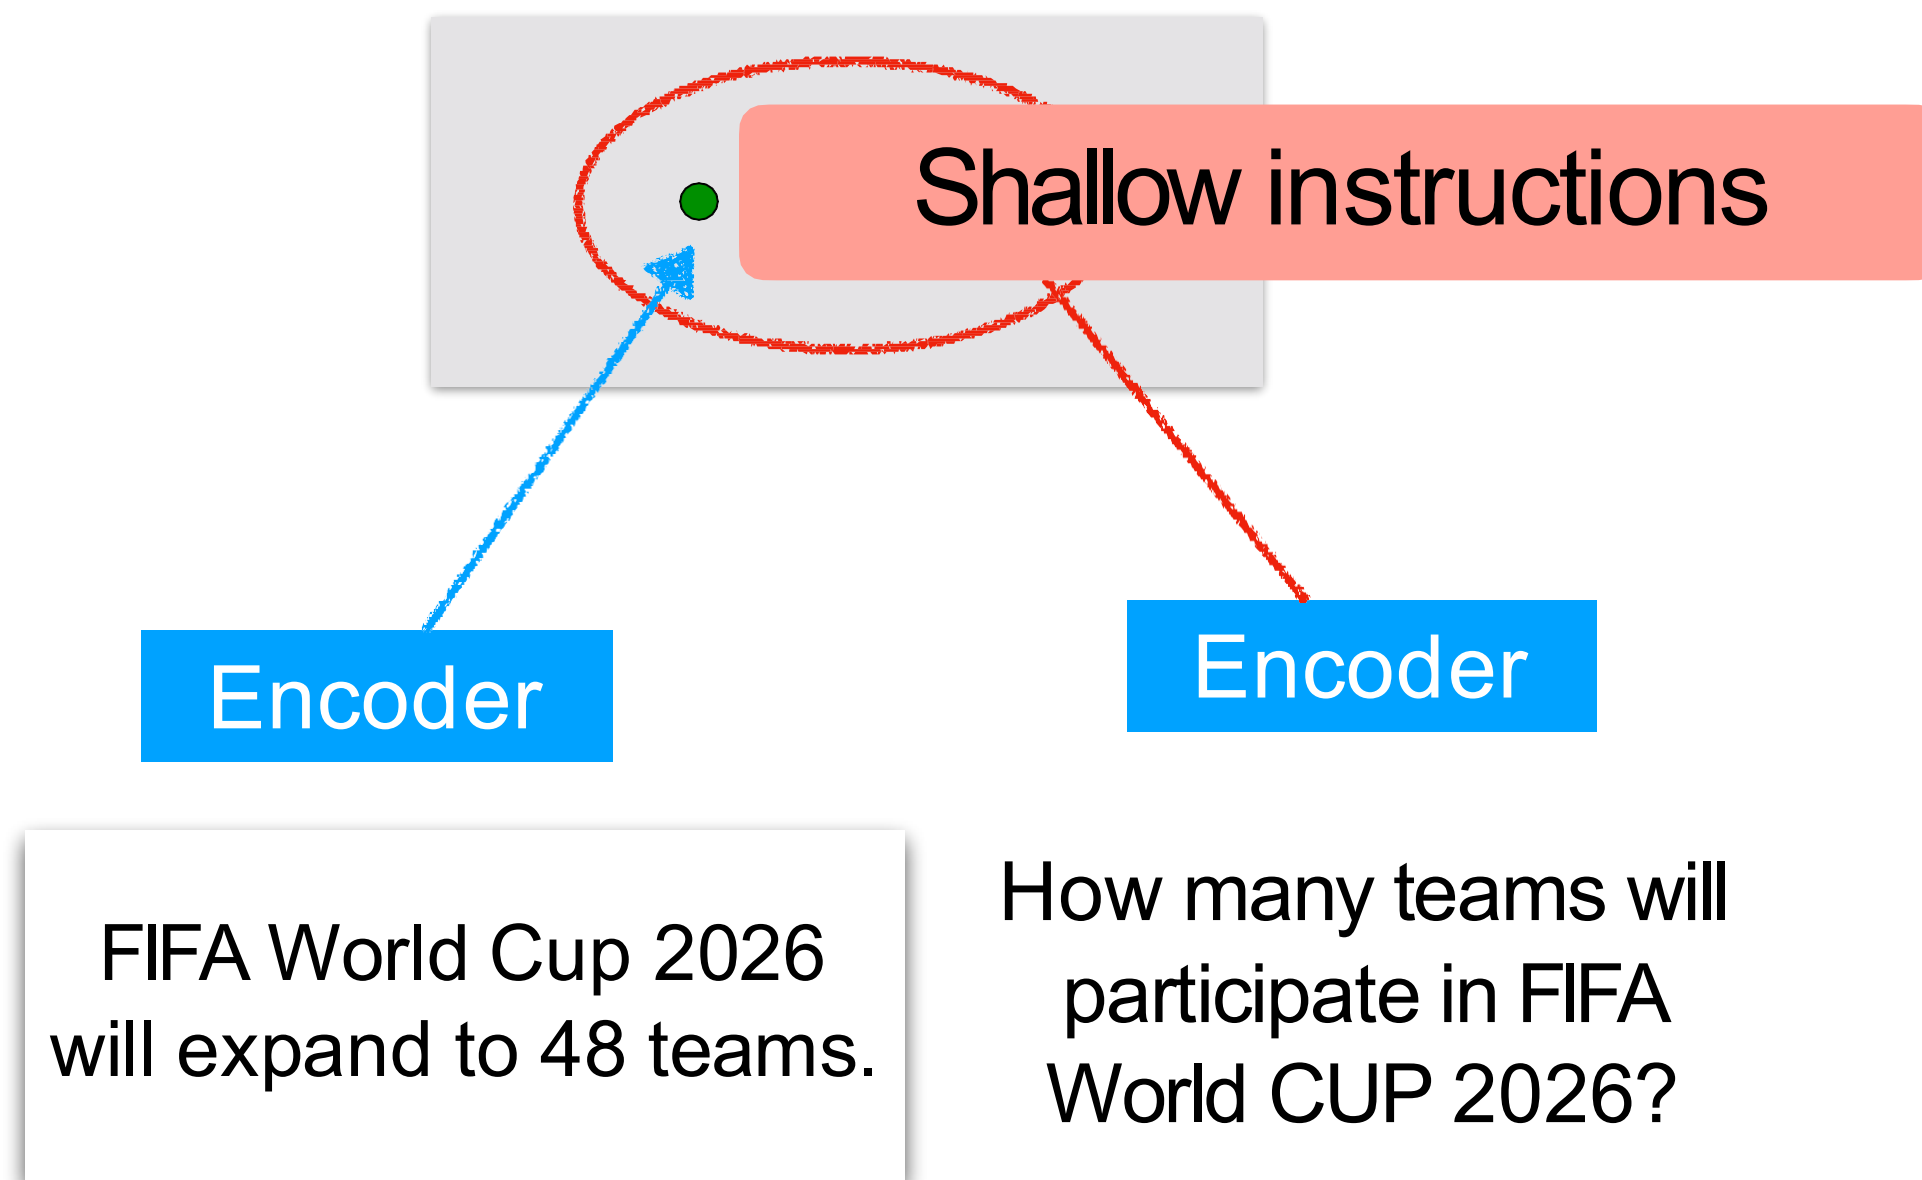
Tasks

Gold documents d^+

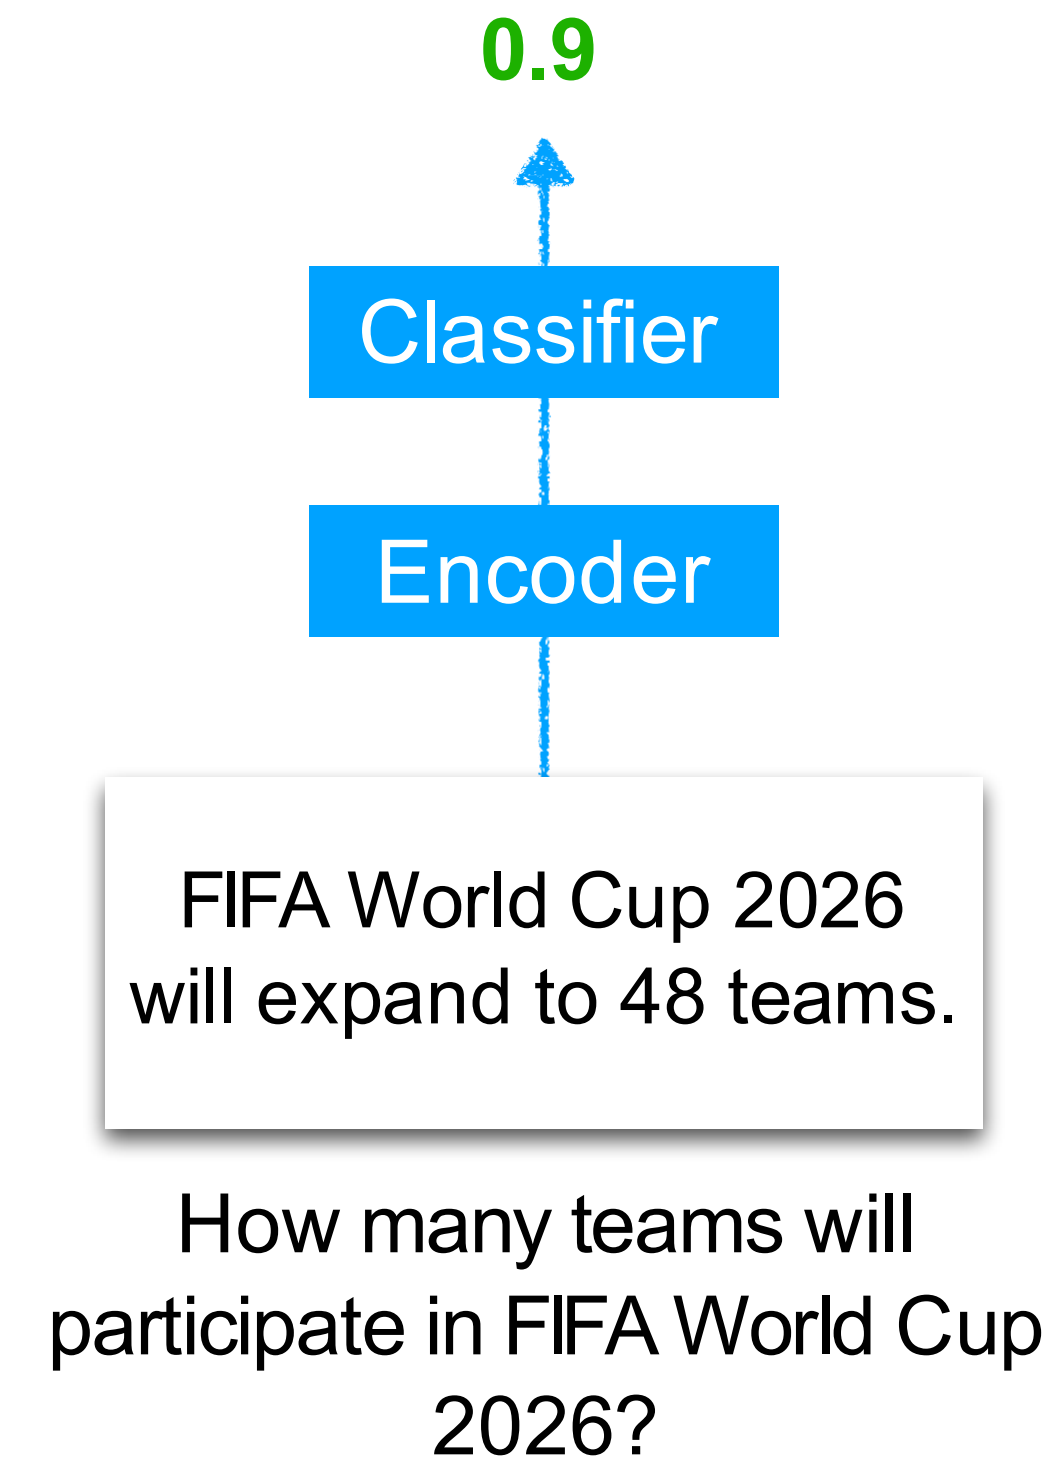
Negative documents d^-

Reranking with Cross-encoder

Bi-Encoder



Cross-Encoder



Evaluation Metrics for Retriever

Evaluation of **unranked** retrieval sets

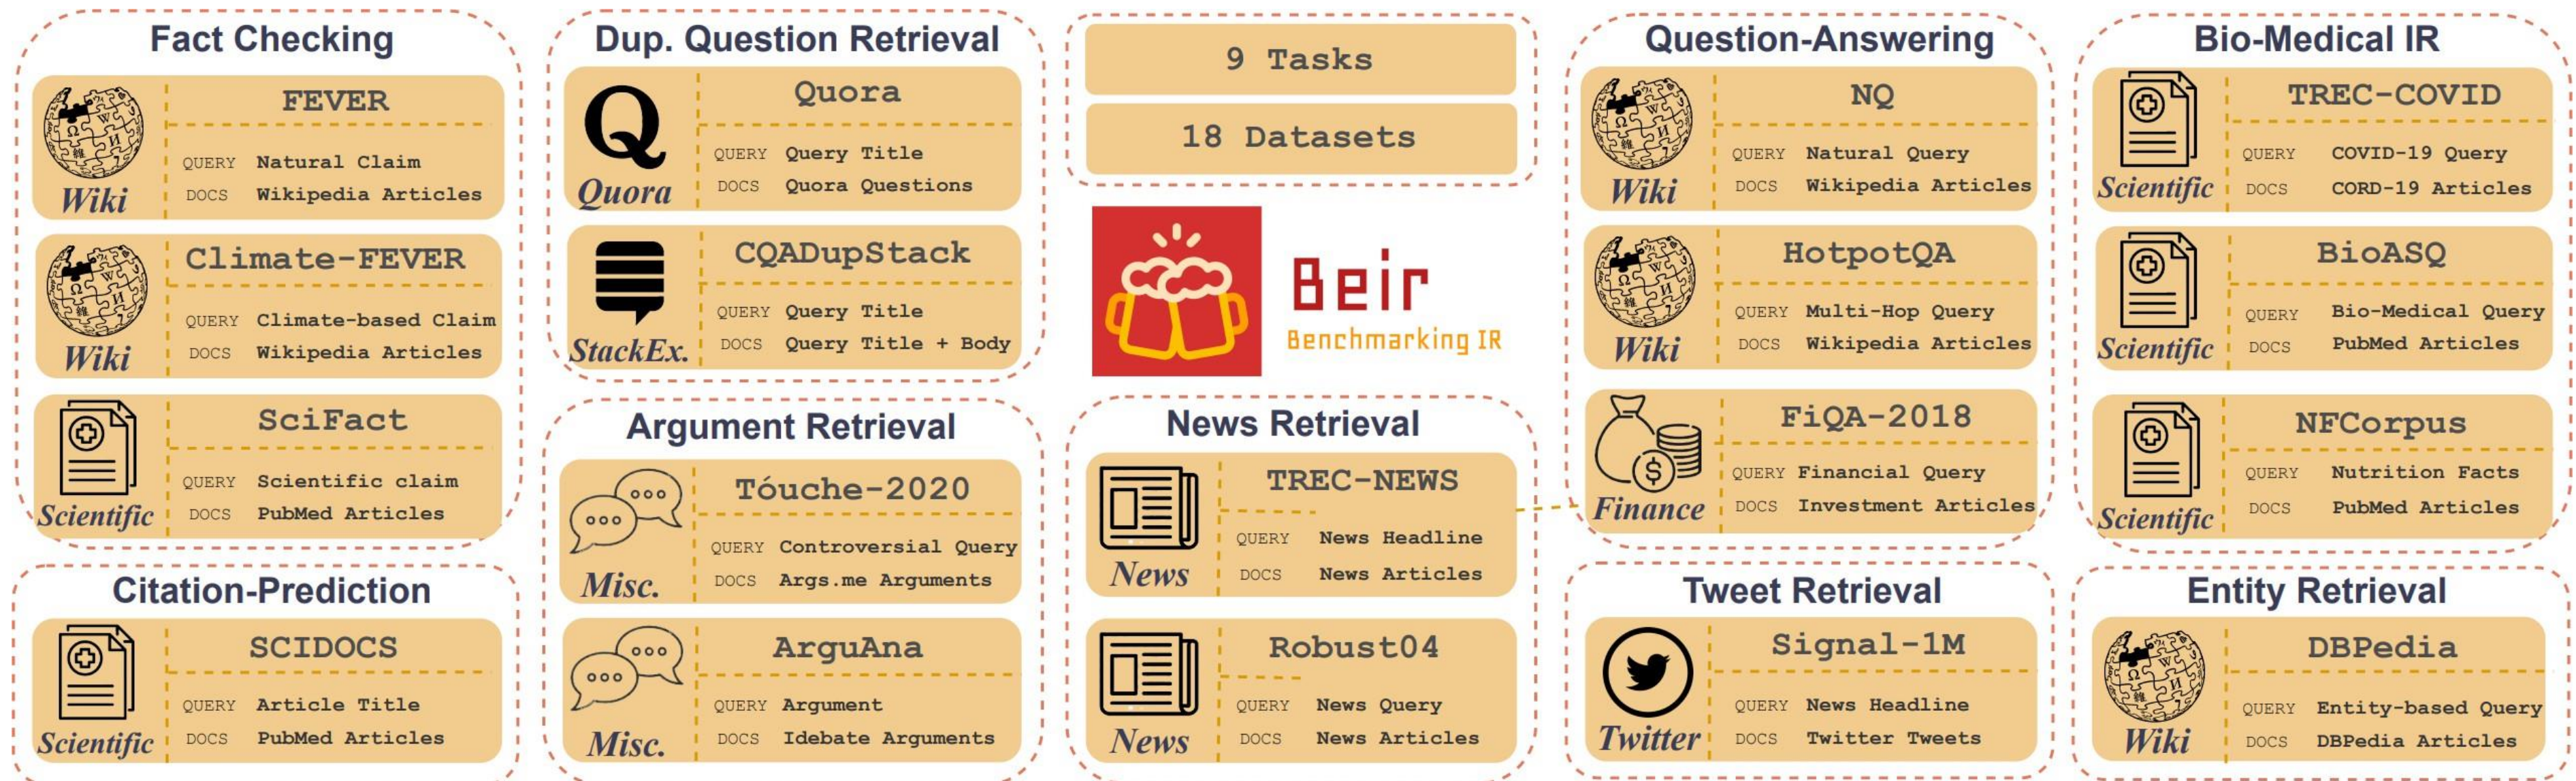
$$\text{Precision} = \frac{\#(\text{relevant items retrieved})}{\#(\text{retrieved items})} \quad \text{Recall} = \frac{\#(\text{relevant items retrieved})}{\#(\text{relevant items})}$$

Evaluation of **ranked** retrieval sets

$$\text{MAP}(Q) = \frac{1}{|Q|} \sum_{j=1}^{|Q|} \frac{1}{m_j} \sum_{k=1}^{m_j} \text{Precision}(R_{jk}) \quad \text{NDCG}(Q, k) = \frac{1}{|Q|} \sum_{j=1}^{|Q|} Z_{kj} \sum_{m=1}^k \frac{2^{R(j,m)} - 1}{\log_2(1 + m)}$$

nDCG@10 is widely used (e.g., BEIR)

Retrieval Benchmarks: BEIR and MTEB



Thakur et al. 2021. NeurIPS D&B. BEIR: A Heterogeneous Benchmark for Zero-shot Evaluation of Information Retrieval Models.

Muennighoff et al. 2022. MTEB: Massive Text Embedding Benchmark.

BEIR Results

	BM25	BM25+CE	DPR	Contriever	
				Ours	Ours+CE
MS MARCO	22.8	41.3	17.7	40.7	47.0
Trec-COVID	65.6	75.7	33.2	59.6	70.1
NFCorpus	32.5	35.0	18.9	32.8	34.4
NQ	32.9	53.3	47.4	49.8	57.7
HotpotQA	60.3	70.7	39.1	63.8	71.5
FiQA	23.6	34.7	11.2	32.9	36.7
ArguAna	31.5	31.1	17.5	44.6	41.3
Touche-2020	36.7	27.1	13.1	23.0	29.8
CQADupStack	29.9	37.0.	15.3	34.5	37.7
Quora	78.9	82.5	24.8	86.5	82.4
DBPedia	31.3	40.9	26.3	41.3	47.1
Scidocs	15.8	16.6	7.7	16.5	17.1
FEVER	75.3	81.9	56.2	75.8	81.9
Climate-FEVER	21.3	25.3	14.8	23.7	25.8
Scifact	66.5	68.8	31.8	67.7	69.2
Avg. w/o CQA	44.0	49.5	26.3	47.5	51.2
Avg.	43.0	48.6	25.5	46.6	50.2
Best on	1	3	0	1	9

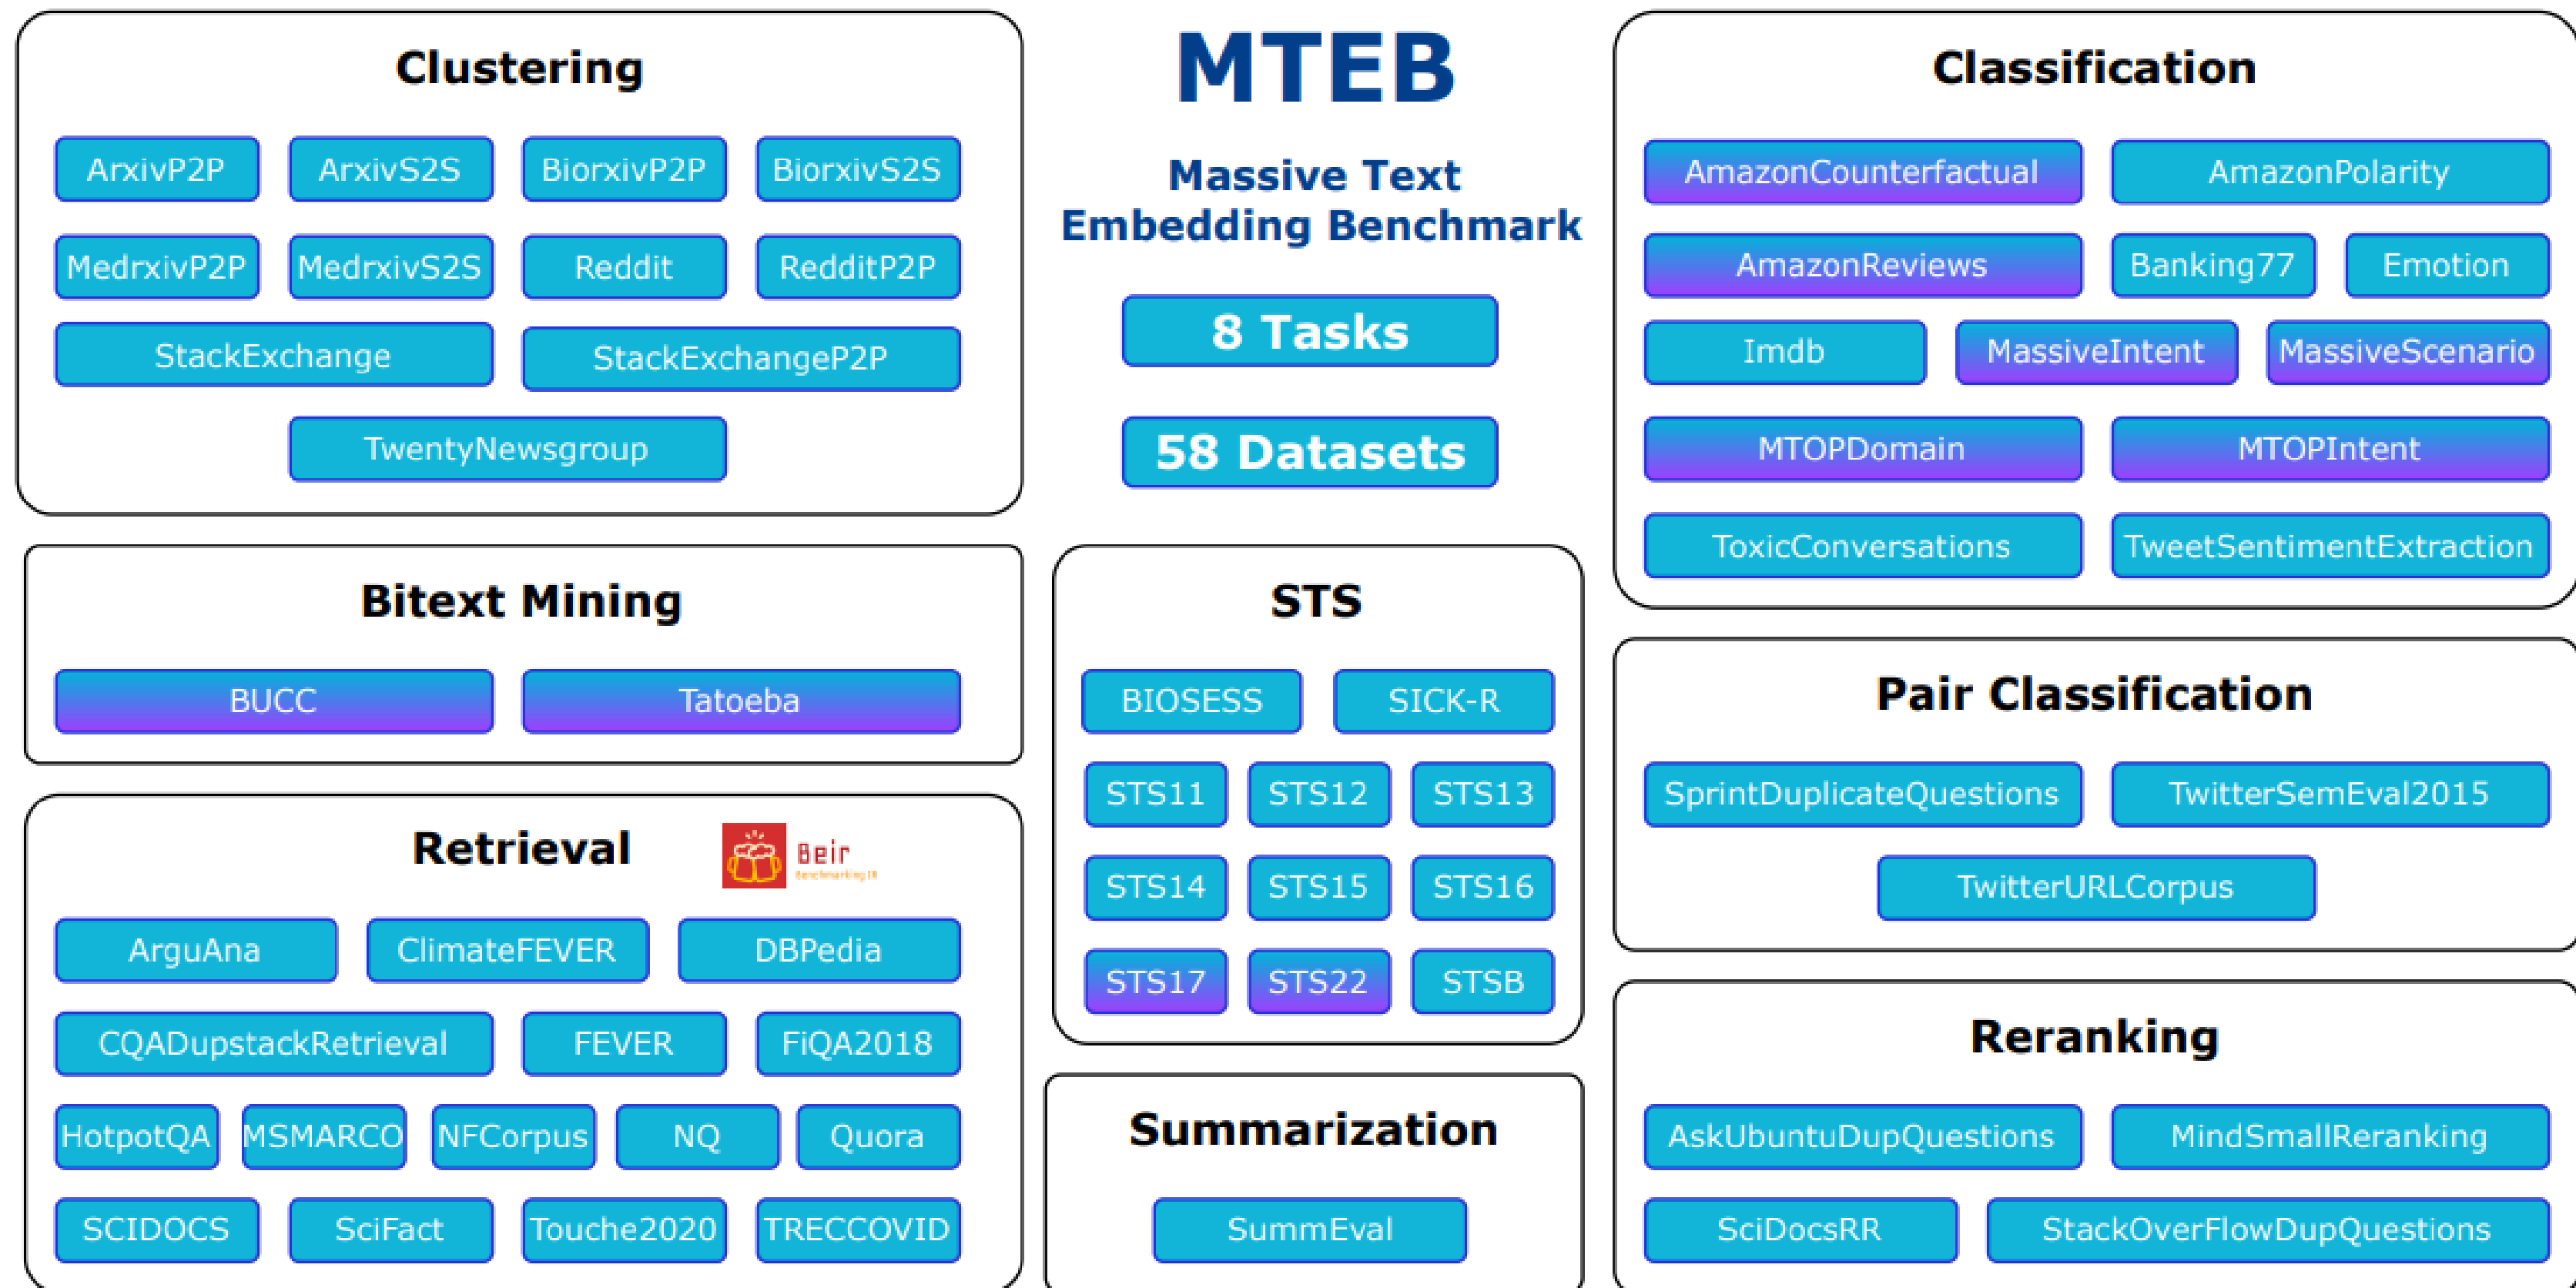
Adding CE (cross-encoder) helps

Dense retrievers could struggle in OOD

Unsupervised training helps in OOD

Izacard et al.TMLR 2022. Unsupervised Dense Information Retrieval with Contrastive Learning.

Retrieval Benchmarks: BEIR and MTEB



Thakur et al. 2021. NeurIPS D&B. BEIR: A Heterogeneous Benchmark for Zero-shot Evaluation of Information Retrieval Models.

MTEB Leaderboard


Rank (Bor...	Model	Zero-shot	Memory U...	Number of P...	Embedding D...	Max Tokens	Mean (T...	Mean (TaskT...	Classificat:
1	QZhou-Embedding	53%	29070	7B	3584	8192	75.97	69.52	88.97
2	LGAI-Embedding-Preview	56%	27125	7B	4096	32768	74.12	68.40	89.97
3	Seed1.5-Embedding	56%	Unknown	Unknown	2048	32768	74.76	68.56	89.88
4	Qwen3-Embedding-8B	95%	28866	7B	4096	32768	75.22	68.71	90.43
5	Seed1.6-embedding	53%	Unknown	Unknown	2048	32768	74.07	67.98	92.42
6	Qwen3-Embedding-4B	95%	15341	4B	2560	32768	74.60	68.10	89.84
7	gemini-embedding-001	95%	Unknown	Unknown	3072	2048	73.30	67.67	90.05
8	jasper en vision language v1	56%	3802	1B	8960	131072	71.41	66.65	90.27
9	Linq-Embed-Mistral	Instruction-tuned retrievers	$q_{\text{inst}}^+ = \text{Instruct} : \{\text{task_definition}\} \text{ Query} : q^+$						
10	SFR-Embedding-Mistral								
11	NV-Embed-v2								
11	NV-Embed-v2	56%	14975	7B	4096	32768	69.81	65.00	87.19

Instruction-tuned retrievers

$q_{\text{inst}}^+ = \text{Instruct} : \{\text{task_definition}\} \text{ Query} : q^+$


<https://huggingface.co/spaces/mteb/leaderboard>

MTEB Leaderboard




<https://huggingface.co/spaces/mteb/leaderboard>

Summary of Part 2




Datastore

Which company developed GPT4o?



Retriever

D




LM

y OpenAI

- ✓ Types of retrievers
- ✓ Training
- ✓ Evaluations

- Different types of retrievers
- Training with contrastive loss
- Common metrics: NDCG@10, Recall ... etc
- Performance v.s. cost trade off

Today's Outline




- ✓ Common architectures
- ✓ Recent progress

Categorizing Retrieval-Augmented LMs

What to retrieve?

Query



Text chunks (passages)?


Tokens?

Something else?

Categorizing Retrieval-Augmented LMs

What to retrieve?

Query



Text chunks (passages)?

Tokens?

Something else?

When to retrieve?

w/ retrieval

The capital city of Ontario is Toronto.

w/ retrieval w/ r w/r w/r w/ r w/r w/r

The capital city of Ontario is Toronto.

w/ retrieval

w/r


w/r

The capital city of Ontario is Toronto.

Categorizing Retrieval-Augmented LMs

What to retrieve?

Query



Text chunks (passages)?

Tokens?

Something else?

When to retrieve?

w/ retrieval

The capital city of Ontario is Toronto.

w/ retrieval w/ r w/r w/r w/ r w/r w/r

The capital city of Ontario is Toronto.

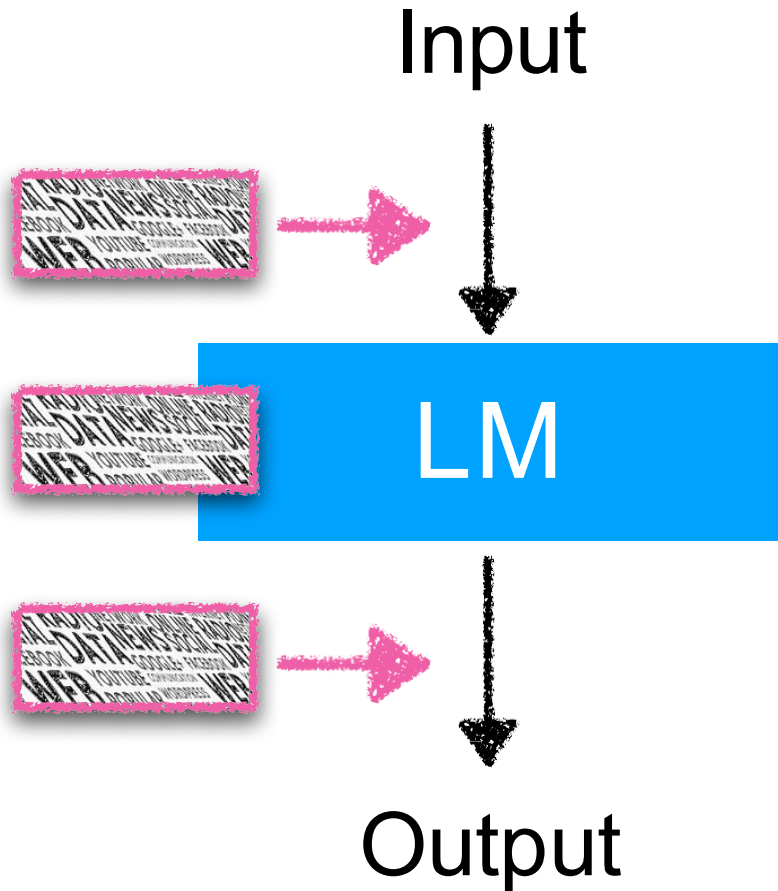
w/ retrieval

w/r

w/r

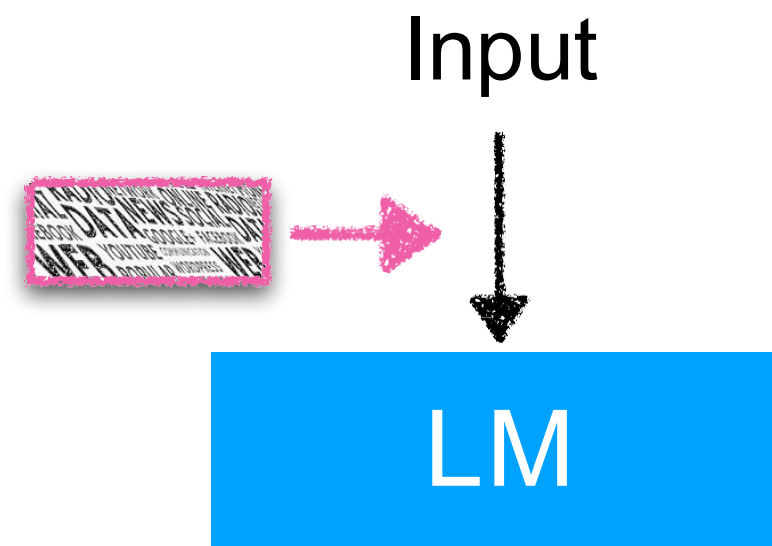
The capital city of Ontario is Toronto.

How to use retrieval?



How to Use Retrieval


Input Augmentation



- Augment input of LMs
- Easy to apply (w/o training) & effective
- Difficulty of using many D

e.g., RAG


Intermediate Fusion



- Modify LMs to incorporate D in intermediate layers
- Scalable to many passages
- Requires retraining

e.g., RETRO, InstructRETRO

Output Interpolation

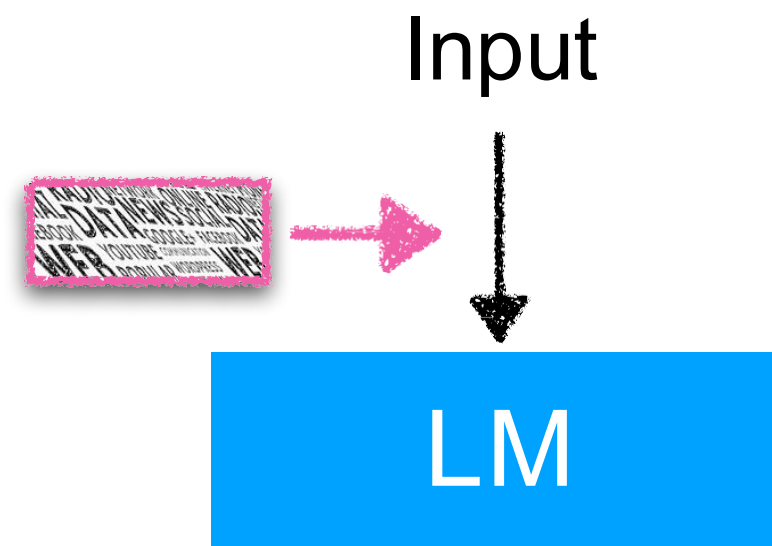


- Directly manipulate output token distributions
- No training required*
- Limited effectiveness on tasks

e.g., kNNLM

How to Use Retrieval

Input Augmentation



- Augment input of LMs
- Easy to apply (w/o training) & effective
- Difficulty of using many D

e.g., RAG

Intermediate Fusion



- Modify LMs to incorporate D in intermediate layers
- Scalable to many passages
- Requires retraining

e.g., RETRO, InstructRETRO


Output Interpolation



- Directly manipulate output token distributions
- No training required*
- Limited effectiveness on tasks

e.g., kNNLM


Retrieval-Augmented Generation (Lewis et al., 2020)



GPT-4o is a pre-trained transformer developed by OpenAI.

GPT4o was released by OpenAI in May 2024.

Retrieval-Augmented Generation (Lewis et al., 2020)




Question: Which company developed GPT4o?

References:

GPT-4o is a pre-trained transformer developed by OpenAI.

GPT4o was released by OpenAI in May 2024.




Training RAG




Independent training

- DPR (Karpukhin et al., 2020)
- DRQA (Chen et al., 2017)



Sequential training


- Evidentiality Generator (Asai et al., 2023)
- REPLUG (Shi et al., 2023)



Joint training

- RAG (Lewis et al., 2021)
- REALM (Guu et al., 2021)

End-to-end training for RAG



$$\sum_j -\log p_{RAG}(y_j | x_j)$$

Minimize NLL as in normal generation training

Retriever score Generator score

$$p_{RAG} \approx \prod_i \sum_{z \in \text{top-}k(p(\cdot|x))} p_\eta(z|x) p_\theta(y_i | x, z, y_{1:i-1})$$


Update retriever encoder and generator

Now people often **combine retrieval with off-the-shelf LMs**

Widely referred to as *RAG*



Effectiveness of Simple RAG

RAG constantly gives performance improvements esp. in long-tail



Limitations of Simple RAG

LMs are easily distracted by unhelpful context



Limitations of Simple RAG



What are the latest discoveries from the James Webb Space Telescope?



The James Webb Space Telescope is designed to peer into the dusty clouds of gas where stars and planetary systems are born. Webb has captured the first direct image of an exoplanet, and the Pillars of Creation in the Eagle Nebula[1][2]. Additionally, the telescope will be used to study the next interstellar interloper[3].

(*Some generated statements may not be fully supported by citations, while others are fully supported.)

Cited Webpages

[1]:  nasa.gov (✗ citation does not support its associated statement)
[NASA's Webb Confirms Its First Exoplanet](#)
... Researchers confirmed an exoplanet, a planet that orbits another star, using NASA's James Webb Space Telescope for the first time. ...

[2]:  cnn.com (⚠ citation partially supports its associated statement)
[Pillars of Creation: James Webb Space Telescope ...](#)
... The Pillars of Creation, in the Eagle Nebula, is a star-forming region captured in a new image (right) by the James Webb Space Telescope that reveals more detail than a 2014 image (left) by Hubble ...

Outputs aren't often supported by citations





Figure 2. Inference time with different input lengths.

Vanilla RAG




Increased latency to encode much longer context


Liu et al. Findings of EMNLP 2023. Evaluating Verifiability in Generative Search Engines

Self-RAG: Learn to Retrieve and Critique


- 🙄 LMs aren't trained with retrieval
- 🙄 Fixed two-stage pipeline




Self-RAG: Learn to Retrieve and Critique



Self-RAG: Learn to Retrieve and Critique




Self-RAG: Learn to Retrieve and Critique



Self-RAG: Learn to Retrieve and Critique

Input

Suggest papers showing LLMs' effectiveness helping scientist to synthesize scientific literature



Sentence 1

Certainly!

Retrieval



Sentence 2

Irrelevant

~~LLMs have been widely used in science.~~

Relevant

Relevant

LLMs have been used in industry widely, such as chatbot system


OpenScholar is a retrieval-augmented LM designed to synthesize literature

GPT4o has shown to be effective to generate new research ideas.

Self-RAG: Learn to Retrieve and Critique

Input

Suggest papers showing LLMs' effectiveness helping scientist to synthesize scientific literature



Sentence 1 Certainly!

Retrieval



Sentence 2

Relevant

Relevant


OpenScholar is a retrieval-augmented LM designed to synthesize literature

GPT4o has shown to be effective to generate new research ideas.

Self-RAG: Learn to Retrieve and Critique

Input

Suggest papers showing LLMs' effectiveness helping scientist to synthesize scientific literature



Sentence 1

Certainly!

Retrieval



Sentence 2

Relevant OpenScholar is an LM for literature synthesis.

Relevant Studies show GPT4o can help scientists for idea generations and literature synthesis.


OpenScholar is a retrieval-augmented LM designed to synthesize literature

GPT4o has shown to be effective to generate new research ideas.

Self-RAG: Learn to Retrieve and Critique

Input

Suggest papers showing LLMs' effectiveness helping scientist to synthesize scientific literature



Sentence 1

Certainly!

Retrieval



Sentence 2

Relevant OpenScholar is an LM for literature synthesis. fully supported

Relevant Studies show GPT4o can help scientists for idea generations and ~~literature synthesis.~~ Partially supported


OpenScholar is a retrieval-augmented LM designed to synthesize literature

GPT4o has shown to be effective to generate new research ideas.

Self-RAG: Learn to Retrieve and Critique

Input

Suggest papers showing LLMs' effectiveness helping scientist to synthesize scientific literature



Sentence 1 Certainly! Retrieval

Sentence 2

Relevant OpenScholar is an LM for literature synthesis. fully supported

0.9


Relevant Studies show GPT4o can help scientists for idea generations and ~~literature synthesis.~~ Partially supported

0.5

Self-RAG: Learn to Retrieve and Critique

Input

Suggest papers showing LLMs' effectiveness helping scientist to synthesize scientific literature



Sentence 1

Certainly!

Retrieval


Sentence 2

Relevant OpenScholar is an LM for
literature synthesis. fully supported

Self-RAG: Learn to Retrieve and Critique

Input

Suggest papers showing LLMs' effectiveness helping scientist to synthesize scientific literature



Sentence 1

Certainly!

Retrieval

Sentence 2


Relevant OpenScholar is an LM for
literature synthesis. fully supported

Sentence 3

Self-RAG: Learn to Retrieve and Critique

Input

Suggest papers showing LLMs' effectiveness helping scientist to synthesize scientific literature



Sentence 1

Certainly!

Retrieval

Sentence 2

Relevant OpenScholar is an LM for
literature synthesis. fully supported


Sentence 3

OpenScholar uses retrieval-augmented
approaches.

Self-RAG: Learn to Retrieve and Critique

Input

Suggest papers showing LLMs' effectiveness helping scientist to synthesize scientific literature



Sentence 1

Certainly!

Retrieval

Sentence 2

Relevant OpenScholar is an LM for
literature synthesis. fully supported

Sentence 3

OpenScholar uses retrieval-augmented
approaches. fully supported

Tool-Augmented LMs

The Brown Act is California's law [WikiSearch("Brown Act") → The Ralph M. Brown Act is an act of the California State Legislature that guarantees the public's right to attend and participate in meetings of local legislative bodies.] that requires legislative bodies, like city councils, to hold their meetings open to the public.

Use search as needed via API calls

Out of 1400 participants, 400 (or [Calculator(400 / 1400) → 0.29] 29%) passed the test.

Can be flexibly extended to other tools

The name derives from "la tortuga", the Spanish word for [MT("tortuga") → turtle] turtle.

Deep Research: Coupling Reasoning LM with Search

Example from WebWalkerQA Dataset
Question: What was the specific schedule for the social event held on the evening after the ACL 2023 best paper awards ceremony? Labeled Answer: The social event was scheduled from 7:00 PM to 10:30 PM on July 11th.
Model Output by WebThinker
Okay, so the user is asking about the specific schedule of a social event that took place on the evening <i>*after*</i> the ACL 2023 best paper awards ceremony. First, I need to figure out when the awards ceremony actually happened. Let me search for that specifically. <code><begin_search_query></code> ACL 2023 best paper awards ceremony date <code><end_search_query></code> <code><begin_search_result></code> The ACL 2023 best paper awards ceremony took place on <i>**July 11, 2023**</i>, according to multiple sources: [...] Web Page 4 [...] explicitly states that the awards were announced on July 11, 2023. [...] Web Page 5 [...] mentions the award presentation occurred on July 11, 2023 [...]. Web Page 6 [...] confirms the awards were announced on July 11, 2023. [...] Web Page 9 [...] also indicates the awards were announced on July 11, 2023. [...] the awards ceremony itself was held on the <i>**11th**</i>. [...] <code><end_search_result></code> Okay, the awards ceremony was definitively on July 11, 2023. The question asks for the social event on the evening <i>*after*</i> the ceremony, which would be the evening of July 12. Let me search for the social events schedule to see what was planned. <code><begin_search_query></code> ACL 2023 social events schedule <code><end_search_query></code> <code><begin_search_result></code> The ACL 2023 social events schedule includes the following relevant details: 1. <i>**Social Event on July 11, 2023**</i>: - <i>**Time**</i>: 7:00 PM – 10:30 PM (EDT) - <i>**Location**</i>: Steam Whistle Brewing [...] - <i>**Description**</i>: This event occurred on the evening of July 11, immediately following the Best Paper Awards ceremony (which took place earlier that day at 16:00–17:30). [...] 2. <i>**No Explicit Social Events on July 12, 2023**</i>: - The evening of July 12, 2023, is not listed as having any official social events in the provided sources. [...] <i>**Note**</i>: The primary social event referenced in the sources is on July 11, not July 12. The Best Paper Awards ceremony was held on July 11, so the subsequent evening event likely refers to the Social Event on that same day. For July 12, only daytime sessions are documented. [...] <code><end_search_result></code> Hmm, interesting. The awards were on July 11, and the search results show a social event on the evening of July 11, immediately following the awards ceremony (7:00 PM - 10:30 PM). However, the results explicitly state there were no official social events scheduled for the evening of July 12, which is technically the evening <i>*after*</i> the July 11 ceremony.

Teach reasoning LMs to plan and search via SFT / RL

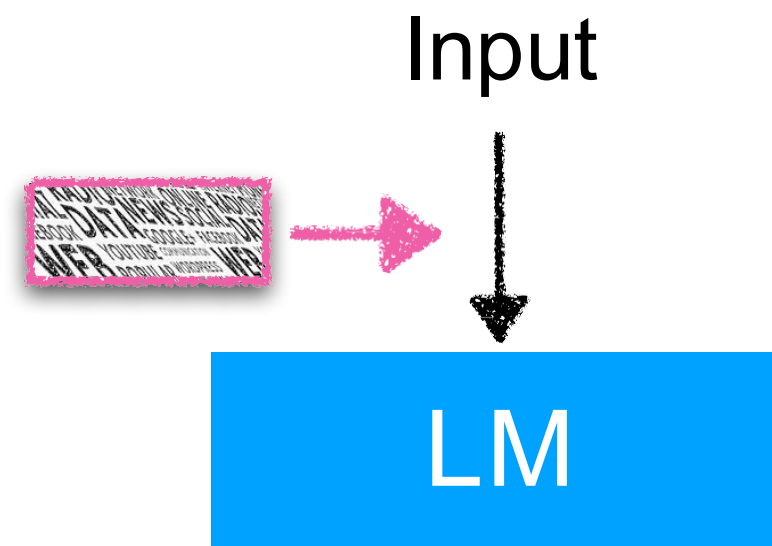
Thinking
Tool call
Search output

Thinking
Tool call
Search output

Thinking

How to Use Retrieval

Input Augmentation



- Augment input of LMs
- Easy to apply (w/o training) & effective
- Difficulty of using many D

e.g., RAG

Intermediate Fusion

Input



Not scalable to many documents (needs context engineering)



- **Not strictly grounded**
- Modifying LMs to incorporate D in intermediate layers
- Scalable to many passages
- Requires retraining

e.g., RETRO, InstructRETRO

Output Interpolation

Output

- Directly manipulate output token distributions
- No training required*
- Limited effectiveness on tasks

e.g., kNNLM

How to Use Retrieval

Input Augmentation

Input



- Augment input of LMs
- Easy to apply (w/o training) & effective
- Difficulty of using many D

e.g., RAG

Intermediate Fusion

Input



- Modify LMs to incorporate D in intermediate layers
- Scalable to many passages
- Requires retraining

e.g., RETRO, InstructRETRO

Output Interpolation



Output



- Directly manipulate output token distributions
- No training required*
- Limited effectiveness on tasks

e.g., kNNLM

RETRO (Borgeaud et al., 2022)



How to Use Retrieval

Input Augmentation

Input



- Augment input of LMs
- Easy to apply (w/o training) & effective
- Difficulty of using many D

e.g., RAG

Intermediate Fusion


Input



- Modify LMs to incorporate D in intermediate layers
- Scalable to many passages
- Requires retraining

e.g., RETRO, InstructRETRO

Output Interpolation




- Directly manipulate output token distributions
- No training required*
- Limited effectiveness on tasks

e.g., kNNLM

kNN-LM (Khandelwal et al. 2020)


Nonparametric distribution



λ : hyperparameter

$$P_{kNN-LM}(y | x) = (1 - \lambda)P_{LM}(y | x) + \lambda P_{kNN}(y | x)$$


Summary of Part 3



- ✓ Common architectures
- ✓ Recent progress

- **RAG** is widely used but several limitations
- Recent progress to overcome such shortcomings
- Other architectures: **intermediate incorporation** or **output interpolation**

Retrieval and Retrieval-Augmented Generation



✓ Sources of datastore

✓ Processing

✓ Scaling

✓ Types of retrievers

✓ Training

✓ Evaluations

✓ Common architectures

✓ Recent progress in RAG

Contact:



<https://akariasai.github.io/>



aasai@andrew.cmu.edu | akaria@openai.com